

Module n°2

Retrieving date

120-007

Auteur : Reda Benkirane
Retrieving date – d September yyyy
Nombre de pages : 33

Summary

1.DISPLAYING DATA FROM MULTIPLE TABLES.....	4
1.1.OBTAINING DATA FROM MULTIPLE TABLES.....	4
1.1.1. <i>Generating Cartesian product.....</i>	4
1.1.2. <i>Types of joins.....</i>	4
1.1.3. <i>Joining tables using Oracle syntax.....</i>	5
1.2. EQUIJOINS.....	5
1.2.1. <i>Retrieving records with equijoins.....</i>	5
1.2.2. <i>Additional search conditions using the AND operator.....</i>	6
1.2.3. <i>Using table aliases.....</i>	7
1.2.4. <i>Joining more than two tables.....</i>	7
1.3. OTHER JOINS.....	8
1.3.1. <i>Retrieving records with Non-Equijoins.....</i>	8
1.3.2. <i>Outer joins.....</i>	9
1.3.3. <i>Self joins.....</i>	10
1.3.4. <i>Cross join.....</i>	10
1.3.5. <i>Natural join.....</i>	11
1.3.6. <i>Creating joins with the USING clause.....</i>	12
1.3.7. <i>Creating joins with the ON clause.....</i>	12
1.3.8. <i>Creating three-way joins with the ON clause.....</i>	13
1.3.9. <i>Left Outer Join and Right Outer Join.....</i>	13
1.3.10. <i>Full Outer Join.....</i>	14
2.AGGREGATING DATA USING GROUP FUNCTIONS.....	15
2.1. GROUP FUNCTIONS.....	15
2.1.1. <i>What are Group functions.....</i>	15
2.1.2. <i>Types of Group Functions.....</i>	16
2.2. USING GROUP FUNCTIONS.....	16
2.2.1. <i>Using the AVG and SUM functions.....</i>	16
2.2.2. <i>Using the MIN and MAX functions.....</i>	16
2.2.3. <i>Using the COUNT function.....</i>	16
2.2.4. <i>Using the DISTINCT keyword.....</i>	17
2.2.5. <i>Using the NVL function with Group Functions.....</i>	17
2.3. CREATING GROUPS OF DATA.....	18
2.3.1. <i>Using the GROUP BY clause.....</i>	18
2.3.2. <i>Using the GROUP BY clause on multiple columns.....</i>	19
2.3.3. <i>Excluding group results using the HAVING clause.....</i>	20
2.3.4. <i>Nesting group functions.....</i>	21
2.3.5. <i>Illegal queries using group functions.....</i>	21
3.SUBQUERIES.....	22
3.1.BASIC SUBQUERIES.....	22
3.1.1. <i>Guidelines for using Subqueries.....</i>	22
3.1.2. <i>Types of subqueries.....</i>	23
3.2.SINGLE-ROW SUBQUERIES.....	23
3.2.1. <i>Executing a Single-Row subquerie.....</i>	23
3.2.2. <i>Using group functions in a subquery.....</i>	24
3.2.3. <i>The HAVING clause with subqueries.....</i>	24
3.3.MULTI-ROW SUBQUERIES.....	25
3.3.1. <i>Using the ANY operator in Multi-Row subqueries.....</i>	25
3.3.2. <i>Using the ALL operator in Multi-Row subqueries.....</i>	25
3.3.3. <i>Null values in a subquery.....</i>	25
3.3.4. <i>Using subqueries in the FROM clause.....</i>	26
4.PRODUCING READABLE OUTPUT WITH ISQL*PLUS.....	27
4.1.SUBSTITUTION VARIABLES.....	27
4.1.1. <i>Using the & substitution variable.....</i>	27
4.1.2. <i>Character and date values with substitution variables</i>	27

4.1.3. Using the <i>&&</i> substitution variable.....	28
4.2. DEFINING SUBSTITUTION VARIABLES.....	28
4.2.1. Using the <i>DEFINE</i> command.....	28
4.2.2. Using then <i>UNDEFINE</i> command.....	29
4.2.3. Using the <i>VERIFY</i> command.....	29
4.3. CUSTOMIZING THE iSQL*PLUS ENVIRONMENT.....	30
4.3.1. Using the <i>SET</i> command.....	30
4.3.2. iSQL*PLUS format commands.....	30
4.3.3. Using the <i>COLUMN</i> command.....	30
4.3.4. Using then <i>BREAK</i> command.....	31
4.3.5. Using the <i>TTITLE</i> and <i>BTITLE</i> commands.....	32
4.3.6. Creating a Script File to run a Report.....	32

1.Displaying data from multiple tables

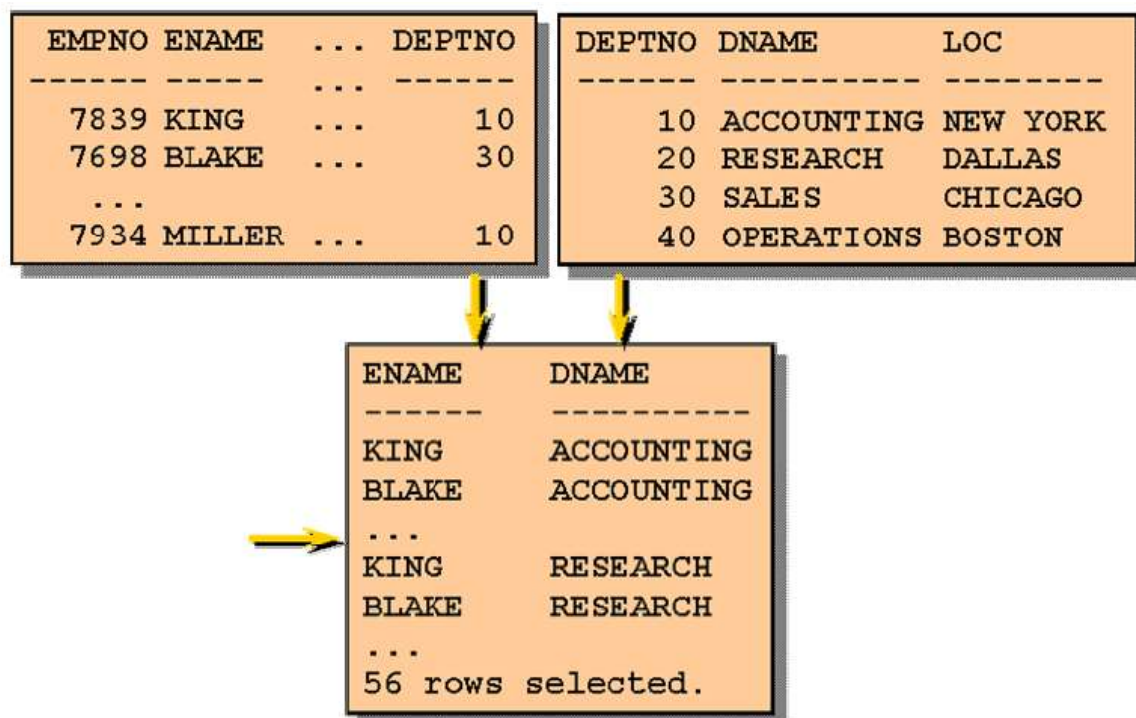
1.1.Obtaining data from multiple tables

1.1.1.Generating Cartesian product

A Cartesian product generates a large number of rows, and the result is not useful.

When a join condition is invalid or omitted, the statement generates a Cartesian product; in which all combinations of rows are displayed.

So, if you don't want to have any Cartesian products, you must include valid join condition in the **WHERE** clause.



The Cartesian product, between the EMP and DEPT table generates 56 rows.

1.1.2.Types of joins

To retrieve information from more than one table in a single SQL statement, you **MUST** join tables using a join condition.

A join condition specifies an existing relation between the columns of these different tables.

Oracle Proprietary Joins (8i and prior)	SQL : 1999 compliant joins
Equijoin	Cross joins
Non-Equijoin	Natural joins
Outer join	Using clause
Self join	Full or two sided outer joins
	Arbitrary join conditions for outer joins

1.1.3. Joining tables using Oracle syntax

To display data from two or more related tables, you just have to write a simple join condition in the **WHERE** clause.

Syntax:

```
SELECT    table1.column, table2.column
FROM      table1, table2
WHERE     table1.column1=table2.column2
```

Table1.column Denotes the table and column from which data is retrieved.

Table1.column1 = Is the condition that joins (or relates) the tables together.

Table2.column2

When joining tables in a SQL statement:

- Precede the columns name with the table name for clarity and performance.
- Must prefix the column name with the table name if the same columns name appears in different tables.
- To join *n* tables, you **MUST** use *n-1* join conditions.

This statements displays employee's department names thanks to a join condition between the EMP and DEPT table.

Values of deptno in the DEPT table are equal to values of deptno on EMP table.

```
SQL> SELECT      emp.ename, dept.dname
2 FROM          emp, dept
3 WHERE         emp.deptno=dept.deptno;

  ENAME      DNAME
-----
SMITH        RESEARCH
ALLEN        SALES
WARD         SALES
JONES        RESEARCH
MARTIN       SALES
BLAKE        SALES
CLARK        ACCOUNTING
SCOTT        RESEARCH
KING         ACCOUNTING
...
15 row(s) selected.
```

You should prefix the name of the columns in the **WHERE** clause with the table name to avoid ambiguity.

For example, without the table prefixes, in the previously example, the *deptno* column could be from the EMP table or the DEPT table.

1.2. Equijoins

1.2.1. Retrieving records with equijoins

An **EQUIJOIN** is used to retrieve information from two or more tables, when there is similarity between values from the different joined tables.

In this example:

- The **SELECT** clause specifies the column names to retrieve:
 - Employee name, employee number and employee department number.
 - Department number and department name.
- The **FROM** clause specifies the two tables that the database must access:
 - EMP table
 - DEPT table
- The **WHERE** clause specifies how the tables are to be joined:
 - EMP.deptno=DEPT.deptno

We use an **EQUIJOIN** because the deptno column is common to the two tables.

```
SQL> SELECT      emp.ename,dept.dname
2 FROM          emp,dept
3 WHERE          emp.deptno=dept.deptno;

  ENAME      DNAME
-----
SMITH        RESEARCH
ALLEN        SALES
WARD         SALES
JONES        RESEARCH
MARTIN       SALES
BLAKE        SALES
CLARK        ACCOUNTING
SCOTT        RESEARCH
KING         ACCOUNTING
...
15 row(s) selected.
```

1.2.2. Additional search conditions using the AND operator

In addition to the join, you can restrict the rows returned using the **AND** operator.

For example, if you want display only the department name of the employee SMITH, you can use the **AND** operator in the **WHERE** clause.

```
SQL> SELECT      emp.ename,dept.dname
  2 FROM          emp,dept
  3 WHERE         emp.deptno=dept.deptno;
  4 AND          emp.ename=' SMITH'
```

ENAME	DNAME
SMITH	RESEARCH

1.2.3.Using table aliases

You can use table aliases instead of table names, if table names are lengthy. Aliases help to keep SQL code smaller, therefore using less memory.

Guidelines:

- Table aliases can be up to 30 characters in length.
- If a table alias is used for a table name in the **FROM** clause, that alias **MUST** be substituted for the table name in the **SELECT** statement.
- Tables' aliases should be meaningful.
- The table alias is valid only for the current **SELECT** statement.
- The aliases are specified in the **FROM** clause after the column name, separated by a space.

Here is an SQL statement joining the DEPT and EMP table and using table's aliases: We use "e" for the EMP table and "d" for the DEPT table.

```
SQL> SELECT      e.ename,d.dname
  2 FROM          emp e,dept d
  3 WHERE         e.deptno=d.deptno;
```

ENAME	DNAME
SMITH	RESEARCH
ALLEN	SALES
WARD	SALES
JONES	RESEARCH
MARTIN	SALES
BLAKE	SALES
CLARK	ACCOUNTING
SCOTT	RESEARCH
KING	ACCOUNTING
...	

15 row(s) selected.

1.2.4.Joining more than two tables

In some cases, you will have to join more than two tables. To create a correct join, you **MUST** create at least $n-1$ join conditions for n tables.

For example, if you want to display the employee name, the department name and the department location, you have to join the EMP table with the DEPT table, and the DEPT table with the LOCATION table.

```
SQL> SELECT      c.name, o.ordid, i.itemid, i.itemtot, o.total
2 FROM          customer c, ord o, item i
3 WHERE         c.custid = o.custid
4 AND           o.ordid = i.ordid
5 AND           c.name = 'TKB SPORT SHOP';
```

NAME	ORDID	ITEMID	ITEMTOT	TOTAL
TKB SPORT SHOP	610	3	58	101.4
TKB SPORT SHOP	610	1	35	101.4
TKB SPORT SHOP	610	2	8.4	101.4

This statement displays orders, item ids, items total and the order total for the customer TBK SPORT SHOP.

1.3. Other joins

1.3.1.Retrieving records with Non-Equijoins

A Non-Equijoin is a join condition containing something other than an equality operator.

EMP

EMPNO	ENAME	SAL
7839	KING	5000
7698	BLAKE	2850
7782	CLARK	2450
7566	JONES	2975
7654	MARTIN	1250
7499	ALLEN	1600
7844	TURNER	1500
7900	JAMES	950
...		

14 rows selected.

SALGRADE

GRADE	LOSAL	HISAL
1	700	1200
2	1201	1400
3	1401	2000
4	2001	3000
5	3001	9999

The salary in the SAL table is enclosed in the LOSAL and HISAL values in the SALGRADE table.

This statement creates a non-equijoin to evaluate an employee's salary grade.
Note that all employees appear once when the statement is executed.


```
SQL> SELECT      e.ename, e.sal, s.grade
  2 FROM          emp e, salgrade s
  3 WHERE         e.sal BETWEEN s.losal AND s.hisal;

  ENAME SAL    GRADE
  -----
  JAMES  950      1
  SMITH  800      1
  ADAMS 1100      1
  ...
14 rows selected.
```

You can use other conditions, such as \geq and \leq ; but **BETWEEN** is still the easiest way.

1.3.2.Outer joins

An outer join is used to see the rows that do not meet the join condition in addition to the rows meeting the condition.

The operator is a plus sign enclosed in parentheses (+), and placed on the side of the join that is deficient in information.

Syntax:

```
SELECT    table1.column, table2.column
FROM      table1, table2
WHERE     table1.column (+) =table2.column;
```

```
SELECT    table1.column, table2.column
FROM      table1, table2
WHERE     table1.column=table2.column (+);
```

Table1.column = Is the condition that joins the tables together.

Table2.column (+) Is the outer join symbol, which can be placed on either side of the **WHERE** clause.

This statement displays the employees' name, with their id and department name.

The OPERATIONS department appears; in spite of the lack of employee in this department.

```
SQL> SELECT      e.ename, d.deptno, d.dname
  2 FROM          emp e, dept d
  3 WHERE         e.deptno (+) = d.deptno
  4 ORDER BY     e.deptno;

  ENAME          DEPTNO DNAME
  -----
  CLARK              10 ACCOUNTING
  KING              10 ACCOUNTING
  MILLER            10 ACCOUNTING
  ...
  TURNER            30  SALES
  WARD              30  SALES
                   40 OPERATIONS
                   50
  ...
```

Note that the outer join operator can appear on only ONE side of the expression, and a condition involving an outer join cannot use the **IN** operator or be linked to another condition by the **OR** operator.

1.3.3. Self joins

A self join is used to join a table to itself.

For example, to find the name of each employee's manager, you must join the EMP table to itself, using a self join.

Syntax:

```
SELECT    alias1.column, alias2.column
FROM      table1 alias1, table2 alias2
WHERE     alias1.column=alias2.column
```

For example, if you want to display the name of each employee's manager, you can execute the following SQL statement:

```
SQL> SELECT      w.ename,m.ename
2  FROM          emp w,emp m
3  WHERE         w.mgr=m.empno;

  ENAME          ENAME
-----
SCOTT            JONES
FORD             JONES
ALLEN            BLAKE
WARD             BLAKE
JAMES            BLAKE
TURNER           BLAKE
...
14 row(s) selected.
```

1.3.4. Cross join

The **CROSS JOIN** returns a Cartesian product from the joined tables.

The result is the same as the Cartesian product between two table.

Cartesian product without cross join:

```
SQL> SELECT      ename, dname
2  FROM          emp,dept;

  ENAME          DNAME
-----
SMITH            ACCOUNTING
ALLEN            ACCOUNTING
WARD             ACCOUNTING
JONES            ACCOUNTING
MARTIN           ACCOUNTING
BLAKE            ACCOUNTING
CLARK            ACCOUNTING
SCOTT            ACCOUNTING
KING             ACCOUNTING
TURNER           ACCOUNTING
ADAMS            ACCOUNTING
...
136 row(s) selected.
```

Cartesian product using cross join:

```
SQL> SELECT      ename, dname
  2  FROM          emp
  3  CROSS JOIN    dept;

  ENAME      DNAME
  -----
SMITH        ACCOUNTING
ALLEN        ACCOUNTING
WARD         ACCOUNTING
JONES        ACCOUNTING
MARTIN       ACCOUNTING
BLAKE        ACCOUNTING
CLARK        ACCOUNTING
SCOTT        ACCOUNTING
KING         ACCOUNTING
TURNER       ACCOUNTING
ADAMS        ACCOUNTING
...
136 row(s) selected.
```

1.3.5.Natural join

The **NATURAL JOIN** is used to let the join be completed automatically by Oracle9i.

The join can happen only on columns having the same names and data types in the two tables.

If the columns have the same name, but different data types, the **NATURAL JOIN** syntax causes an error.

The **NATURAL JOIN** will produce the same result as an equijoin.

Equijoin using the **WHERE** clause:

```
SQL> SELECT      emp.ename, dept.dname
  2  FROM          emp, dept
  3  WHERE          emp.deptno=dept.deptno;

  ENAME      DNAME
  -----
SMITH        RESEARCH
ALLEN        SALES
WARD         SALES
JONES        RESEARCH
MARTIN       SALES
BLAKE        SALES
CLARK        ACCOUNTING
SCOTT        RESEARCH
KING         ACCOUNTING
...
15 row(s) selected.
```

Equijoin using the **NATURAL JOIN** clause:

```
SQL> SELECT      ename, dname
2  FROM          emp
3  NATURAL JOIN dept;
```

ENAME	DNAME
SMITH	RESEARCH
ALLEN	SALES
WARD	SALES
JONES	RESEARCH
MARTIN	SALES
BLAKE	SALES
CLARK	ACCOUNTING
SCOTT	RESEARCH
KING	ACCOUNTING
...	

15 row(s) selected.

1.3.6. Creating joins with the USING clause

If columns have the same name but the data types do not match, the **NATURAL JOIN** clause fails, but it can be modified with the **USING** clause to specify the columns that should be used for an equijoin.

It's important not to use table name or alias in the referenced columns.

```
SQL> SELECT      e.empno,e.ename,d.loc
2  FROM          emp e JOIN dept d
3  USING          (deptno);
```

EMPNO	ENAME	LOC
7369	SMITH	DALLAS
7499	ALLEN	CHICAGO
7521	WARD	CHICAGO
7566	JONES	DALLAS
7654	MARTIN	CHICAGO
7698	BLAKE	CHICAGO
7782	CLARK	NEW YORK
7788	SCOTT	DALLAS
...		

15 row(s) selected.

1.3.7. Creating joins with the ON clause

The join condition for the **NATURAL** join is basically an equijoin of all columns with the same name, but you can specify arbitrary conditions or specify columns to join using the **ON** clause.

```
SQL> SELECT      e.empno,e.ename,d.loc,d.dname
  2 FROM          emp e JOIN dept d
  3 ON            (e.deptno=d.deptno);
```

EMPNO	ENAME	LOC	DNAME
7369	SMITH	DALLAS	RESEARCH
7499	ALLEN	CHICAGO	SALES
7521	WARD	CHICAGO	SALES
7566	JONES	DALLAS	RESEARCH
7654	MARTIN	CHICAGO	SALES
7698	BLAKE	CHICAGO	SALES
7782	CLARK	NEW YORK	ACCOUNTING
7788	SCOTT	DALLAS	RESEARCH
...			

15 row(s) selected.

Note that the **ON** clause can also be used as follows to join columns that have different names.

```
SQL> SELECT      e.ename emp, r.mgr mgr
  2 FROM          emp e JOIN emp r
  3 ON            (e.mgr = r.empno);
```

EMP	MGR
SCOTT	7839
FORD	7839
ALLEN	7839
WARD	7839
JAMES	7839
...	

14 row(s) selected.

1.3.8. Creating three-way joins with the ON clause

A three-way join is a join of three tables. Joins are performed from left to right.

Using the **ON** clause with the **JOIN** clause, you can join three tables

```
SQL> SELECT      c.name,o.ordid,i.itemid,i.itemtot,o.total
  2 FROM          customer c
  3 JOIN          ord o
  4 ON            (c.custid=o.custid)
  5 JOIN          item t
  6 ON            (o.ordid=i.ordid);
```

NAME	ORDID	ITEMID	ITEMTOT	TOTAL
TKB SPORT SHOP	610	3	58	101.4
TKB SPORT SHOP	610	1	35	101.4
TKB SPORT SHOP	610	2	8.4	101.4

3 rows selected.

1.3.9. Left Outer Join and Right Outer Join

This join returns the same results as the return of the classical outer join, the only difference is just that instead of putting the plus sign on the left of the **WHERE** clause, you have to specify the **RIGHT** keyword.

This SQL statement displays all rows in the EMP table, which is the right table even if there is no match in the DEPT table.

```
SQL> SELECT      e.ename, d.deptno, d.dname
2  FROM          emp e
3  RIGTH OUTER JOIN dept d
4  ON            e.deptno=d.deptno);
```

ENAME	DEPTNO	DNAME
CLARK	10	ACCOUNTING
KING	10	ACCOUNTING
MILLER	10	ACCOUNTING
...		
TURNER	30	SALES
WARD	30	SALES
	40	OPERATIONS
	50	
...		

1.3.10.Full Outer Join

The **FULL OUTER JOIN** gives all rows from the tables even if there is not match between the joined tables. You can apply additional conditions in the **WHERE** clause using the **AND** operator.

Syntax:

```
SELECT          table1.column1, table2.column
FROM            table1
FULL OUTER JOIN dept
ON( table1.column2=table2.column2).
```

```
SQL> SELECT      e.ename,e.deptno,d.dname
2  FROM          emp e
3  FULL OUTER JOIN dept d
4  ON            (e.deptno=d.deptno);
```

ENAME	DEPTNO	DNAME
MILLER	10	ACCOUNTING
KING	10	ACCOUNTING
CLARK	10	ACCOUNTING
...		
PAPIER	30	SALES
JAMES	30	SALES
TURNER	30	SALES
...		
WARD	30	SALES
ALLEN	30	SALES
THG		
THG		
		store
		STORE
		OPERATIONS
		STORE

22 row(s) selected.

2. Aggregating Data using Group Functions

2.1. Group functions

2.1.1. What are Group functions

EMPLOYEES

DEPARTMENT_ID	SALARY
90	24000
90	17000
90	17000
60	9000
60	6000
60	4200
50	5800
50	3500
50	3100
50	2600
50	2500
110	12000
110	8300

20 rows selected.

The maximum salary in the EMPLOYEES table.

MAX(SALARY)
24000

Group functions operate on sets of rows to give one result per group. These sets of rows may be the whole table or the table split into smaller groups.

Syntax:

```

SELECT      [column,] group_function (column) , ...
FROM        table
[WHERE]      condition]
[GROUP BY]  column]
[ORDER BY]  column]
  
```

2.1.2.Types of Group Functions

Here are group functions whose accept arguments:

Function	Description
AVG ([DISTINCT ALL] <i>n</i>)	Average value of <i>n</i> , ignoring null values.
COUNT ({* [DISTINCT ALL] expr {})	Numbers of rows, where <i>expr</i> evaluates to something other than null.
MAX ([DISTINCT ALL] <i>expr</i> {})	Maximum value of <i>expr</i> , ignoring null values.
MIN ([DISTINCT ALL] <i>expr</i> {})	Minimum value of <i>expr</i> , ignoring null values.
STDDEV ([DISTINCT ALL] <i>x</i>)	Standard deviation of <i>x</i> , ignoring null values
SUM ([DISTINCT ALL] <i>x</i>)	Sum values of <i>x</i> , ignoring null values.
VARIANCE ([DISTINCT ALL] <i>x</i>)	Variance of <i>x</i> , ignoring null values.

2.2. Using Group Functions

2.2.1.Using the AVG and SUM functions

You can use **AVG**, **SUM**, **MIN** and **MAX** functions with columns storing numeric data.

This SQL statement displays, the average, highest, lowest, and sum of salaries of all the employees.

SQL> SELECT	AVG (sal) , MAX (sal) , MIN (sal) , SUM (sal)			
2 FROM	emp;			
	AVG (SAL)	MAX (SAL)	MIN (SAL)	SUM (SAL)
	-----	-----	-----	-----
	2001,66667	5000	800	30025

2.2.2.Using the MIN and MAX functions

The **MIN** and **MAX** functions can be used with any datatype.

This statement displays the highest salary and the employee name that is last in the alphabetized list of all employees.

SQL> SELECT	MAX (SAL) , MAX (ename)	
2 FROM	emp;	
	MAX (SAL)	MAX (ENAME)
	-----	-----
	5000	WARD

Notice that the **AVG**, **SUM**, **VARIANCE** and **STDDEV** functions can be used only with numeric datatypes.

2.2.3.Using the COUNT function

The **COUNT (*)** returns the number of rows in a table that satisfy the criteria of the **SELECT** statement, including duplicate rows and null values.

If a **WHERE** clause is included in the **SELECT** statement, the **COUNT (*)** returns number of rows that satisfy the condition.

COUNT(*expr*) returns the number of non-null values in the *expr* column.

COUNT(DISTINCT *expr*) returns the number of unique, non-null values in the *expr* column.

This SQL statement displays the number of all rows in the EMP table:

```
SQL> SELECT      COUNT ( *)
  2 FROM          emp;

COUNT ( *)
-----
          17
```

This statement displays, the number of employees working in the department 30:

```
SQL> SELECT      COUNT (ename)
  2 FROM          emp
  3 WHERE         deptno=30;

COUNT (ENAME)
-----
              7
```

2.2.4.Using the DISTINCT keyword

You can use the **DISTINCT** keyword to suppress the counting of any duplicate values within a column.

This example displays the number of distinct department values in the EMP table:

```
SQL> SELECT      COUNT (DISTINCT deptno)
  2 FROM          emp;

COUNT (DISTINCTDEPTNO)
-----
              3
```

2.2.5.Using the NVL function with Group Functions

By default, group functions ignore null values, but using the **NVL** function, it's possible to force group functions to include null values.

If you need to calculate the annual salary of all employees, you have to include their commission, but some employees don't have commission.

Using the **NVL** function, the average is calculated even if the commission is null.

```
SQL> SELECT      SAL*12+NVL(comm,0) as "SALARY"
2  FROM          emp;

SALARY
-----
12600
19700
15500
35700
16400
34200
29400
```

2.3. Creating groups of data

2.3.1. Using the GROUP BY clause

You can divide the table of information into smaller groups, using the **GROUP BY** clause.

Syntax:

```
SELECT      column, group_function(column)
FROM        table
[WHERE      condition]
[GROUP BY expression]
[ORDER BY column];
```

Guideline:

- When including a group function in a **SELECT** statement, you can't select individual results as well, UNLESS the individual column appears in the **GROUP BY** clause.
- You can exclude rows before dividing them into groups, using a **WHERE** clause.
- You must include the columns that are not used with the group functions in the **SELECT** clause into the **GROUP BY** clause.
- You can't use column alias in the **GROUP BY** clause.
- By default, rows are sorted by ascending order.

When using the **GROUP BY** clause, you have to make sure that all columns in the **SELECT** list, which are not used in the group functions, are included in the **GROUP BY** clause.

This SQL statement displays the number of employees working in the different departments:

```
SQL> SELECT      COUNT(ename),deptno
2 FROM          emp
3 GROUP BY      deptno;
```

COUNT (ENAME)	DEPTNO
3	10
5	20
7	30

Note that columns do not have to be in the SELECT list to be used is the GROUP BY clause.

```
SQL> SELECT      COUNT(ename)
2 FROM          emp
3 GROUP BY      deptno;
```

COUNT (ENAME)
3
5
7

2.3.2.Using the GROUP BY clause on multiple columns

Sometimes, you may need to see results for groups within groups.

You can return summary results for groups and subgroups by listing more than one column in the GROUP BY clause.

EMPLOYEES

DEPARTMENT_ID	JOB_ID	SALARY
10	AD_ASSI	4400
20	MK_MAN	13000
20	MK_REP	6000
50	ST_CLERK	3500
50	ST_CLERK	3100
50	ST_CLERK	2600
50	ST_CLERK	2500
50	ST_MAN	5800
60	IT_PROG	9000
60	IT_PROG	6000
60	IT_PROG	4200
80	SA_MAN	10500
80	SA_REP	11000
110	AC_MGR	12000
	SA_REP	7000

20 rows selected.

“Add up the salaries in the EMPLOYEES table for each job, grouped by department.”

DEPARTMENT_ID	JOB_ID	SUM(SALARY)
10	AD_ASST	4400
20	MK_MAN	13000
20	MK_REP	6000
50	ST_CLERK	11700
50	ST_MAN	5800
60	IT_PROG	19200
80	SA_MAN	10500
80	SA_REP	19600
90	AD_PRES	24000
90	AD_VP	34000
110	AC_ACCOUNT	8300
110	AC_MGR	12000
	SA_REP	7000

13 rows selected.

In this SQL statement:

- The SELECT clause specifies the columns to retrieve.
- The FROM clause specifies tables that database must access.
- The GROUP BY clause specifies how results must be grouped:
 - First, the rows are grouped by department number.

- Second, within the department number groups, the rows are grouped by job.

```
SQL> SELECT      deptno, job, SUM(sal)
  2 FROM          emp
  3 GROUP BY      deptno, job;

  DEPTNO JOB          SUM(SAL)
-----
      10 CLERK          1300
      10 MANAGER        2450
      10 PRESIDENT      5000
      20 ANALYST        6000
      20 CLERK          1900
      20 MANAGER        2975
      30 CLERK           950
      30 MANAGER        2850
      30 SALESMAN       6600

      10 row(s) selected.
```

2.3.3. Excluding group results using the HAVING clause

If you want to restrict groups of rows that you select, you have to use the HAVING clause.

Syntax:

```
SELECT column, group_function
FROM table
[WHERE condition]
[GROUP BY group_by_expression]
[HAVING group_condition]
[ORDER BY column];
```

When using the **HAVING** clause, the Oracle Server performs the following steps:

1. Rows are grouped.
2. The group function is applied to the group.
3. The groups that match the criteria in the **HAVING** clause are displayed.

It's recommended to place the **HAVING** clause after the **GROUP BY** clause.

This SQL statement displays department id and maximum salaries for those departments whose maximum salary is greater or equal to 3000\$.

```
SQL> SELECT      MAX(sal),deptno
2 FROM          emp
3 GROUP BY      deptno
4 HAVING        MAX(sal)>=3000;
```

MAX (SAL)	DEPTNO
5000	10
3000	20

2.3.4.Nesting group functions

If needed, group functions can be nested to a depth of two.
The following statement displays the maximum average salary.

```
SQL> SELECT      MAX (AVG (sal))
2 FROM          emp
3 GROUP BY      deptno;
```

MAX (AVG (SAL))
2916,66667

2.3.5.Illegal queries using group functions

Any column or expression in the **SELECT** list that is not an aggregate function must be in the **GROUP BY** clause.

```
SQL> SELECT      deptno,COUNT(ename)
2 FROM          emp;
SELECT deptno,COUNT(ename)
*
```

ERROR at line 1 :
ORA-00937: Not a single-group group function

You cannot use the **WHERE** to restrict groups, and group functions cannot be used in the **WHERE** clause.

```
SQL> SELECT      deptno,AVG(sal)
2 FROM          emp
3 WHERE          AVG(sal)=1000;
WHERE AVG(sal)=1000
*
```

ERROR at line 1 :
ORA-00934: Group function is not allowed here

Column alias cannot be used in the **GROUP BY** clause.

```
SQL> SELECT      deptno department,AVG(SAL)
2 FROM          emp
3 GROUP BY      department;
GROUP BY department
*
```

ERROR at line 1 :
ORA-00904: Invalid column name

3.Subqueries

3.1.Basic Subqueries

3.1.1.Guidelines for using Subqueries

A subquery is a **SELECT** statement that is embedded in a clause of another **SELECT** statement. Using subqueries, you can build powerful statements out of simple ones.

Subqueries can be placed in the **WHERE** clause, the **HAVING** clause and the **FROM** clause. You can include to your statements operators to add comparison conditions such as '=' or **IN**.

In the following statement, the inner query determines the salary of SMITH, and the outer query takes the result of the inner one and uses this result to display all employee that earn more than this amount.

```
SQL> SELECT      emp.ename
2 FROM          emp
3 WHERE         sal > (SELECT sal
4 FROM          emp
5 WHERE         ename='SMITH') ;

      ENAME
-----
      ALLEN
      WARD
      JONES
      MARTIN
      ...
14 row(s) selected.
```

Guidelines:

- Enclose subqueries in parentheses.
- Place subqueries on the right side of the comparison condition.
- The **ORDER BY** clause in the subquery is not needed unless you are performing Top-N analysis.
- Use single-row operator with single-row subqueries and use multiple-row operators with multiple-row subqueries.

3.1.2.Types of subqueries

- **Single-row subquery**



- **Multiple-row subquery**



Type	Description
Single-row	Queries that return only one row from the inner SELECT statement.
Multiple-row	Queries that return more than one row from the inner SELECT statement.

3.2.Single-Row subqueries

3.2.1.Executing a Single-Row subquerie

A single row subquery is one that returns one row from the inner **SELECT** statement. This type of subquery uses single-row operators, and are placed in the **WHERE** clause.

This statement displays employees that work in the same department as the employee number 7369.

```
SQL> SELECT      ename
2 FROM          emp
3 WHERE         deptno=(SELECT deptno
4                   FROM emp
5                   WHERE empno=7369) ;

      ENAME
-----
      SMITH
      JONES
      SCOTT
      ADAMS
      FORD
```

Many subqueries can be included into a statement using operators like **AND** and **OR**.

This following statement displays employees working in the same department as the employee SMITH, and has a salary greater than the salary of the employee number 7369.

```

SQL> SELECT      ename
2  FROM          emp
3  WHERE          deptno=(SELECT deptno
4                      FROM emp
5                      WHERE ename='SMITH')
6  AND           sal>
7                      (SELECT sal
8                      FROM emp
9                      WHERE empno=7369);

      ENAME
-----
      JONES
      SCOTT
      ADAMS
      FORD

```

3.2.2.Using group functions in a subquery

Group functions can be used in single-row subqueries.

The following SQL statement displays the name, the job, and the salary of employees that have a salary equal to the lowest salary.

```

SQL> SELECT      ename,job,sal
2  FROM          emp
3  WHERE          sal=(SELECT MIN(sal)
4                      FROM emp);

      ENAME      JOB      SAL
-----
      SMITH      CLERK      800

```

3.2.3.The HAVING clause with subqueries

As the **WHERE** clause, the **HAVING** clause can contain subqueries.

This statement displays departments that have lowest salary greater than the lowest salary of the department 20.

```

SQL> SELECT      deptno,MIN(sal)
2  FROM          emp
3  GROUP BY      deptno
4  HAVING         MIN(SAL)>
5                      (SELECT MIN(SAL)
6                      FROM emp
7                      WHERE deptno=20);

      DEPTNO      MIN (SAL)
-----
          10          1300
          30          950

```


This SQL statement displays the job which has the lowest average salary.

```
SQL> SELECT      job,AVG (SAL)
  2  FROM          emp
  3  GROUP BY      job
  4  HAVING         AVG (SAL) =  (SELECT MIN (AVG (sal))
  5                               FROM emp
  6                               GROUP BY job);
```

JOB	AVG (SAL)
CLERK	1037,5

3.3.Multi-Row subqueries

3.3.1.Using the ANY operator in Multi-Row subqueries

The **ANY** operator compares a specified value to each values returned by the Multiple-row subquery.

Comparison operator	Meaning
>ANY	More than the minimum
>ALL	More than the maximum
<ANY	Less than the maximum
<ALL	Less than the minimum

This SQL statement displays employees who are not CLERKS and whose salary is less than the salary of any CLERKS.

```
SQL> SELECT      empno,ename,job,sal
  2  FROM          emp
  3  WHERE          sal< ANY (SELECT sal
  4                      FROM emp
  5                      WHERE job='CLERK')
  6  AND            job <> 'CLERK';
```

EMPNO	ENAME	JOB	SAL
7521	WARD	SALESMAN	1250
7654	MARTIN	SALESMAN	1250
7952	PAPIER	SALESMAN	1000

3.3.2.Using the ALL operator in Multi-Row subqueries

The **ALL** operator compares a value to every value returned by a subquery.
Notice that the **NOT** operator can be used with **IN**, **ANY**, and **ALL** operators.

The following statement displays all employees whose salary is less than the salary of employees in the department 30.

```
SQL> SELECT      ename
  2  FROM          emp
  3  WHERE          sal<  ALL (SELECT  sal  FROM  emp  WHERE
deptno=30);
```

ENAME
SMITH

3.3.3.Null values in a subquery

If the subquery returns a null value, the outer query will not return any rows.

If you want, the statement to return results, you have to use the **NVL** function.

3.3.4. Using subqueries in the FROM clause

Subqueries can be written in the **FROM** clause.

The subquery's result is a virtual table, which must have an alias.

Example:

SQL>	SELECT	e.ename, e.sal, e.deptno, b.salavg
2	FROM	emp e, (SELECT deptno, AVG(sal) salavg
3		FROM emp
4		GROUP BY deptno) b
5	WHERE	e.deptno = b.deptno
6	AND	e.sal < b.salavg;
ENAME	SAL	DEPTNO SALAVG

CLARK	2450	10 2916,66667
MILLER	1300	10 2916,66667
SMITH	800	20 2175
ADAMS	1100	20 2175
PAPIER	1000	30 1485,71429
MARTIN	1250	30 1485,71429
JAMES	950	30 1485,71429
WARD	1250	30 1485,71429
8 row(s) selected.		

This SQL statement displays the name, the salary, the department number and the department salary average for every employee earning a salary less than the average salary of his department.

The subquery returns department numbers and average salary of every department, and these results are stored in the virtual table.

The EMP table and the virtual table *b* are joined.

4. Producing readable output with iSQL*PLUS

4.1. Substitution variables

4.1.1. Using the & substitution variable

In iSQL*PLUS, you can use a single ampersand (&) substitution variable to temporarily store values. You can predefine variables in iSQL*PLUS by using the **DEFINE** command, which creates and assigns a value to a variable.

Use an ampersand to identify variables in your SQL statement. You don't need to specify a value for each variable.

This statement displays the department name, and the department location of the department number given by the user. (10)

```
SQL> SELECT      dname, loc
   2  FROM        dept
   3  WHERE      deptno= &deptid;
Enter a value for deptid : 10
old   3 : WHERE deptno= &deptid
new   3 : WHERE deptno= 10

DNAME          LOC
-----
ACCOUNTING     NEW YORK
```

4.1.2. Character and date values with substitution variables

When using substitution variables with dates or characters, you must enclose values within single quotation marks so that when the user enters the value, he doesn't have to use single quotation marks.

This SQL statement displays employees whose job is the same as the one specified by the user

```
SQL> SELECT      ename
  2 FROM          emp
  3 WHERE         job='&job';
Enter a value for job : CLERK
old   3 : WHERE job='&job'
new   3 : WHERE job='CLERK'

ENAME
-----
SMITH
ADAMS
...
4 row(s) selected.
```

Note that all values are case sensitive.

4.1.3.Using the && substitution variable

If a variable is preceded by a double ampersand (&&), iSQL*PLUS will prompt the user for its value only one time.

The user has to enter only one time a value for the *column_name* variable, which appears twice in the statement.

The value entered by the user, is used twice.

```
SQL> SELECT      ename,job,&&column_n
  2 FROM          emp
  3 ORDER BY      &&column_name
  4 ;

Enter a column for column_na

old   1 : SELECT ename,job,&&
new   1 : SELECT ename,job,s
old   3 : ORDER BY &&column_n
new   3 : ORDER BY sal

ENAME      JOB      SAL
-----
SMITH      CLERK      800
JAMES      CLERK      950
PAPIER     SALESMAN    1000
ADAMS      CLERK      1100
...
17 row(s) selected.
```

4.2.Defining substitution variables

4.2.1.Using the DEFINE command

Command	Description
DEFINE variable=variable	Creates a user variable with the CHAR data type and assigns it a value.
DEFINE variable	Displays the variable, its value, and its data type
DEFINE	Displays all user variables with their values and data types.

The first **DEFINE** command defines the *deptname* variable and attributes it a value.

The second **DEFINE** command displays the value of the deptname variable.

The SQL statement displays information about the department which is specified by the deptname variable.

```
SQL> DEFINE deptname=sales
SQL> DEFINE deptname
DEFINE DEPTNAME          = "sales" (CHAR)
SQL> SELECT
2  FROM          dept
3  WHERE          dname=UPPER('&deptname');

old  3 : WHERE dname=UPPER('&deptname')
new  3 : WHERE dname=UPPER('sales')

-----
DEPTNO DNAME          LOC
-----
      30 SALES          CHICAGO
```

4.2.2.Using then UNDEFINE command

The **UNDEFINE** command is used to delete a user variable.

```
SQL> UNDEFINE deptname
SQL> DEFINE deptname
SP2-0135: the symbol deptname is UNDEFINED
```

4.2.3.Using the VERIFY command

To confirm the changes in the SQL statement, you can use the **iSQL*PLUS VERIFY** command. Setting **VERIFY ON**, forces iSQL*PLUS to display the text of a command before and after substitution variables are replaced with values.

Entrez les intructions :

```
SET VERIFY ON
SELECT empno,ename
FROM emp
WHERE empno=&empid;
```

ancien 3 : WHERE empno=&empid
nouveau 3 : WHERE empno=7369

EMPNO	ENAME
7369	SMITH

4.3. Customizing the iSQL*PLUS environment

4.3.1. Using the SET command

SET command variables

SET variable and values	Description
ARRAY [SIZE] {20 n}	Sets the database data fetch size.
FEED [BACK] {6 n OFF ON}	Displays the number of records returned by a query when the query selects at least <i>n</i> records.
HEA[DING] {OFF ON}	Determines whether column headings are displayed in reports.
LONG {80 n}	Sets the maximum width for displaying LONG values.

The **SET** commands are used to control current session, so the values are available only for the current session and will not be saved.

Syntax:

SET *system_variable value*

In the syntax:

System_variable Is a variable that controls one aspect of the session environment.

Value Is a value for the system variable.

It is possible to verify what has been set, using the **SHOW** command.
To see all **SET** variable values, use the **SHOW ALL** command.

4.3.2. iSQL*PLUS format commands

If you want to control the report features, you have to use the following commands:

Command	Description
COL[UMN] [column option]	Controls column format
TTI[TLE] [text OFF ON]	Specifies a header to appear at the top of each pages of the report.
BTI[TLE] [text OFF ON]	Specifies a footer to appear at the bottom of each pages of the report.
BRE[AK] [ON <i>report_element</i>]	Suppresses duplicate values and divides rows of data into sections by using line breaks.

Format commands remain effective until the end of the iSQL*PLUS session or until the format setting is overwritten or cleared.

4.3.3. Using the COLUMN command

This command controls the display of column.

COLUMN command options:

Option	Description
CLE[AR]	Clears any column formats
HEA[DING] <i>text</i>	Sets the column heading (a vertical line forces a line feed in the heading if you do not use justification)
FOR[MAT] <i>format</i>	Changes the display of the column data
NOPRI[NT]	Hides the column
NUL[L] <i>text</i>	Specifies the <i>text</i> to be displayed for null values
PRI[NT]	Shows the column

COLUMN format models:

Element	Description	Example	Result
9	Single Zero-suppression digit	999999	1234
0	Enforces leading zero	099999	001234
\$	Floating dollar sign	\$9999	\$1234
L	Local currency	L9999	L1234
.	Position of decimal point	9999.99	1234.00
,	Thousand separator	9,999	1,234

Example:

- Create column headings:

```
o SQL> COLUMN ename HEADING 'Employee|Name'
o SQL> COLUMN sal JUSTIFY LEFT FORMAT $99,990.00
o SQL> COLUMN mgr FORMAT 999999999 NULL 'No manager'
```

- Displays the current setting for the ENAME column:

```
o SQL> COLUMN ename
o COLUMN      ename ON
o HEADING     'Employee|Name' headsep '|' '
```

- Clear settings for the ENAME column:

```
SQL> COLUMN ename CLEAR
```

4.3.4.Using then BREAK command

You can use the **BREAK** command to divide rows into sections and suppress duplicate values. Use the **ORDER BY** clause that you are breaking to ensure that the **BREAK** command works.

Syntax:**BREAK** *on column* [| *alias* | *row*]

column [| *alias* | *row*] Suppresses the display of duplicate values for a given column.

You can clear all **BREAK** settings by using the **CLEAR** command:

CLEAR BREAK**4.3.5.Using the TTITLE and BTITLE commands**

TTITLE is used to format page headers and **BTITLE** for footers, which appear at the bottom of the page.

The **BTITLE** and **TTITLE** commands have the same syntax.

Syntax:**TTI[TLE] | BTI[TLE]** [*text* | **OFF** | **ON**]

text: Represents the title text.
OFF | ON: Toggles the title either off or on. It's not visible when turned off.

Example:

* Set the report header:

```
SQL> TTITLE 'Salary|Report'
```

* Set the report footer:

```
SQL> BTITLE 'Confidential'
```

4.3.6.Creating a Script File to run a Report**How to create a script file:**

1. Create the SQL **SELECT** statement at the SQL prompt. You must ensure that the data required for the report is accurate before you save the statement to a file and apply formatting commands.
2. Save the **SELECT** statement to a script file.
3. Edit the script file to enter iSQL*PLUS commands.
4. Add the required formatting commands before the **SELECT** statement. Be sure not to place iSQL*PLUS commands within the **SELECT** statement.
5. Verify that the **SELECT** statement is followed by a run character, either a semicolon (;) or a slash (/).

6. Add the format-creating iSQL*PLUS commands after the run character.
7. Save the script file with your changes.
8. Load the script file into the iSQL*PLUS text window, and click the Execute button.

This script displays the job, name, and salary for every employee whose salary is less than \$3000. The script adds a centred, two-line header that reads “Employee Report” and a centred footer that reads “Confidential”.

It renames the job title column to read “Job name” splits over two lines, renames the employee column name to read “Employee” and finally, renames the salary column to read “Salary” and formats it as \$2,500.00.

Note that REM represents a remark or comment in iSQL*PLUS.

```
SET FEEDBACK OFF
TTITLE 'Employee|Report'
BTITLE 'Confidential'
BREAK ON job
COLUMN job HEADING 'Job|Name'
COLUMN ename HEADING 'Employee'
COLUMN sal HEADING 'Salary' FORMAT $99,999.99
REM ** Insert SELECT statement
SELECT job,ename,sal
FROM emp
WHERE sal < 3000
ORDER BY job,ename
/
REM clear all formatting commands ...
SET FEEDBACK ON
COLUMN job CLEAR
COLUMN ename CLEAR
COLUMN sal CLEAR
CLEAR BREAK
```