

Module n°4

Advanced SQL

1Z0-007

Auteur : Reda Benkirane
Advanced SQL – d September yyyy
Nombre de pages : 47


Summary

1.USING SET OPERATORS.....	4
1.1.THE SET OPERATORS.....	4
1.2.UNION AND UNION ALL.....	5
1.2.1.The UNION operator.....	5
1.2.2.Using the UNION operator.....	5
1.2.3.The UNION ALL operator.....	6
1.2.4.Using the UNION ALL operator.....	6
1.3.INTERSECT	7
1.3.1.The INTERSECT operator.....	7
1.3.2.Using the INTERSECT operator.....	7
1.4.MINUS.....	8
1.4.1.The MINUS operator.....	8
1.4.2.Using the MINUS operator.....	8
1.5.SET OPERATOR GUIDELINES.....	9
1.5.1.SET operators rules.....	9
1.5.2.Matching the SELECT statement.....	9
1.6.THE ORACLE SERVER AND SET OPERATORS.....	9
1.7.CONTROLLING THE ORDER OF ROWS.....	10
2.ORACLE SERVER DATETIME FUNCTIONS.....	11
2.1.DATETIME FUNCTIONS.....	11
2.1.1.TZ_OFFSET.....	11
2.1.2.CURRENT_DATE.....	11
2.1.3.CURRENT_TIMESTAMP.....	12
2.1.4.LOCALTIMESTAMP.....	13
2.1.5.DBTIMEZONE and SESSIONTIMEZONE.....	14
2.1.6.EXTRACT.....	14
2.2.CONVERSIONS.....	16
2.2.1.TIMESTAMP conversion using FROM_TZ.....	16
2.2.2.STRING to TIMESTAMP using TO_TIMESTAMP and TO_TIMESTAMP_TZ.....	17
2.2.3.Time interval conversion with TO_YMINTERVAL.....	17
3.GROUP BY WITH ROLLUP AND CUBE OPERATORS.....	19
3.1.REVIEWS OF GROUP BY AND HAVING.....	19
3.1.1.Review of group functions.....	19
3.1.2.Review of the GROUP BY clause.....	19
3.1.3.Review of the HAVING clause.....	19
3.2.ROLLUP.....	19
3.2.1.The ROLLUP operator.....	19
3.2.2.Example.....	19
3.3.CUBE.....	20
3.3.1.The CUBE operator.....	20
3.3.2.Example.....	20
3.4. ANALYTICAL FUNCTIONS.....	21
3.4.1.Describe analytical functions.....	21
3.4.2.RANK function.....	21
3.4.3.CUME_DIST function.....	22
3.5.GROUPING.....	23
3.5.1.The GROUPING function.....	23
3.5.2.Example.....	23
3.5.3.GROUPING SETS.....	24
3.5.4.Using GROUPING SETS.....	25
3.6.COMPOSITE COLUMNS.....	25
3.6.1.The composite columns.....	25
3.6.2.Example.....	26
3.7.CONCATENATED GROUPINGS.....	27
3.7.1.The concatenated groupings.....	27

3.7.2.Example.....	27
4.ADVANCED SUBQUERIES.....	28
4.1.SUBQUERIES.....	28
4.1.1.What is a subquery.....	28
4.1.2.Multiple-column subqueries	28
4.1.3.Using subqueries.....	28
4.2.COLUMN COMPARISONS.....	28
4.2.1.Pairwise comparisons.....	28
4.2.2.Nonpairwise comparisons.....	29
4.3.SCALAR SUBQUERIES.....	29
4.3.1.The scalar subqueries.....	29
4.3.2.Example.....	30
4.4.CORRELATED SUBQUERIES.....	30
4.4.1.The correlated subqueries.....	30
4.4.2.Example.....	31
4.5.EXISTS AND NOT EXISTS OPERATORS.....	31
4.5.1.The EXISTS operator.....	31
4.5.2.Using the EXISTS operator.....	31
4.5.3.The NOT EXISTS operator.....	32
4.6.CORRELATED UPDATE AND DELETE.....	33
4.6.1.The correlated UPDATE.....	33
4.6.2.The correlated DELETE.....	34
4.7.WITH CLAUSE.....	34
4.7.1.The WITH clause.....	34
4.7.2.Example.....	34
5.HIERARCHICAL RETRIEVAL.....	36
5.1. HIERARCHICAL QUERIES OVERVIEW.....	36
5.1.1.When is a hierarchical query possible?.....	36
5.1.2.Natural structure tree.....	36
5.1.3.Hierarchical queries.....	36
5.2. WALKING THE TREE.....	36
5.2.1.Start point.....	36
5.2.2.Direction.....	37
5.2.3.Example.....	37
5.3. ORGANIZING DATA.....	37
5.3.1.Ranking rows with the LEVEL pseudo column.....	37
5.3.2.Formatting hierarchical report using LEVEL and LPAD.....	38
5.3.3.Pruning branches.....	38
5.3.4.Ordering Data.....	39
5.3.5.ROW_NUMBER() function.....	39
6.DML AND DDL STATEMENTS.....	41
6.1.MULTITABLE INSERT STATEMENT.....	41
6.1.1.Types of multitable INSERT statement.....	41
6.1.2.Unconditional INSERT ALL.....	42
6.1.3.Conditional INSERT ALL.....	42
6.1.4.Conditional FIRST INSERT.....	43
6.1.5.Pivoting INSERT.....	43
6.2.EXTERNAL TABLES.....	44
6.2.1.Creating an external table.....	44
6.2.2.Example of creating an external table.....	44
6.2.3.Querying external tables.....	46
6.3.CREATE INDEX WITH CREATE TABLE STATEMENT.....	46

1.Using SET operators

1.1.The SET operators



The set operators combine the results of two or more component queries into one result. Queries containing set operators are called *compound queries*.

Operator	Result
INTERSECT	All distinct rows selected by both queries (INTERSECT combines two queries and returns only those rows from the first SELECT statements that are identical to at least one row from the second SELECT statement.)
UNION	All distinct rows selected by both query
UNION ALL	All rows selected by both query, including all duplicates
MINUS	All distinct rows selected by the first SELECT statement that are not produced in the second SELECT statement


```
SELECT      column1, column2, ...
FROM        table1
SET OPERATOR
SELECT      column1, column2, ...
FROM table2 ;
```

All set operators have equal precedence. If a SQL statement contains multiple set operators, the database evaluates them from left (top) to right (bottom) if no parentheses explicitly specify another order.

INTERSECT and **MINUS** operators are specific to Oracle.

1.2.UNION and UNION ALL

1.2.1.The UNION operator



The **UNION** operator combines results from multiple queries, but eliminates duplicate rows. Use the **UNION** operator to return all rows from multiple tables and eliminate any duplicate rows.

The number of columns and the data types of those columns must be identical in the two **SELECT** statements. The names of the columns need not be identical.

UNION operates over all of the columns being selected.

NULL values are not ignored during duplicate checking.

Queries that use **UNION** in the **WHERE** clause must have the same number and type of columns in their **SELECT** list.

By default, the output is sorted in ascending order of the first column of the **SELECT** clause.

1.2.2.Using the UNION operator

The **UNION** operator eliminates all duplicate rows.

Example:


```
SQL> SELECT      ename, job, sal
  2  FROM      emp
  3  UNION
  4  SELECT      name, title, 0
  5  FROM      emp_history;
```

ENAME	JOB	SAL
ADAMS	CLERK	1100
ALLEN	SALESMAN	0
ALLEN	SALESMAN	1600
BALFORD	CLERK	0
BLAKE	MANAGER	2850
...		

23 row(s) selected

→ This statement displays all the results from the two queries, but common lines are not displayed. Also it displays a 0 when the line comes from the table that doesn't have a column for the salaries.

1.2.3.The UNION ALL operator



Use the **UNION ALL** operator to return all rows from multiple queries.
Unlike **UNION**, duplicate rows are not eliminated and the output is not sorted by default.
The **DISTINCT** keyword cannot be used

1.2.4.Using the UNION ALL operator

The **UNION ALL** operator returns results from both queries, including duplicate rows.

Example:


```
SQL> SELECT      ename, empno, job
  2  FROM          emp
  3  UNION ALL
  4  SELECT      name, empid, title
  5  FROM          emp_history
  6  ORDER BY      ename;
```

```
ENAME          EMPNO JOB
-----
ADAMS           7876 CLERK
ALLEN           7499 SALESMAN
ALLEN           7499 SALESMAN
BALFORD         6235 CLERK
BLAKE           7698 MANAGER
BRIGGS          7225 PAY CLERK
...
23 row(s) selected.
```

→ This statement displays all the results from the two queries, including the ones that are on the two tables.

1.3.INTERSECT

1.3.1.The INTERSECT operator



Use the **INTERSECT** operator to return all rows common to both queries.

The number of columns and the data types of those columns must be identical in the two **SELECT** statements. The names of the columns need not be identical.

Reversing the order of the intersected tables does not alter the result.

INTERSECT, like **UNION**, does not ignore **NULL** values.

Queries that use **INTERSECT** in the **WHERE** clause must have the same number and type of columns in their **SELECT** list.

1.3.2.Using the INTERSECT operator

Example:

```
SQL> SELECT      ename, empno, job
2  FROM          emp
3  INTERSECT
4  SELECT      name, empid, title
5  FROM          emp_history;
```


ENAME	EMPNO	JOB
-----	-----	-----
ALLEN	7499	SALESMAN
CLARK	7782	MANAGER
SCOTT	7788	ANALYST

→ This statement returns only the records that have the same values in the selected columns in both tables.

If you add a column in the **SELECT** list, the result may be different.

1.4.MINUS

1.4.1.The MINUS operator



Use the **MINUS** operator to return rows returned by the first query but not the second query (the first **SELECT** statement **MINUS** the second **SELECT** statement).

Like other operators, you must pay attention to the number of columns and their data type. Beware of the **WHERE** clause too.

1.4.2.Using the MINUS operator

Example:

```
SQL> SELECT      name, empid, title
  2  FROM          emp_history
  3  MINUS
  4  SELECT      ename, empno, job
  5  FROM          emp;

NAME          EMPID TITLE
-----
BALFORD        6235 CLERK
BRIGGS         7225 PAY CLERK
JEWELL         7001 ANALYST
SPENCER        6087 OPERATOR
VANDYKE        6185 MANAGER
WILD           7356 DIRECTOR

6 row(s) selected.
```

→ This statement returns the name, the employee id and their function that have quit the company (the employees from the EMP_HISTORY table minus the employees from the EMP table).

1.5.SET operator guidelines

1.5.1.SET operators rules

The expressions in the **SELECT** lists must match in number and data type. The names displayed in the result are the ones from the first **SELECT** list.

Duplicate rows are automatically eliminated except in **UNION ALL**.

The output is sorted in ascending order by default except in **UNION ALL**. The **ORDER BY** clause can be used, but only in the last query. The argument of the **ORDER BY** clause must match one of the names in the first **SELECT** list.

Column names from the first query appear in the result.

Parentheses can be used to alter the sequence of execution.

Queries that use set operator in their **WHERE** clause must have the same number and type of columns in their **SELECT** list.

1.5.2. Matching the SELECT statement

Because the expressions in the **SELECT** lists of compound queries must match in number and data type, you can use dummy columns and the data type conversion functions to comply with this rule. This can be done with the DUAL table and with various functions of data converting.

Example:

```
SQL> SELECT      deptno, TO_CHAR(null) AS location, hiredate
  2 FROM          emp
  3 UNION
  4 SELECT      deptno, loc, TO_DATE(null)
  5 FROM          dept;
```

DEPTNO	LOCATION	HIREDATE
10	NEW YORK	
10		09/06/81
10		17/11/81
10		23/01/82
20	DALLAS	
20		17/12/80
...		
30		28/09/81
30		03/12/81
DEPTNO	LOCATION	HIREDATE
40	BOSTON	

18 row(s) selected.

→ This statement displays deptno, location and hiredate of all employees.

1.6.The Oracle Server and SET operators

When using **SET** operators, the Oracle Server eliminates automatically duplicate rows, except when using the **UNION ALL** operator.

By default, the result is sorted in ascending order of the first column of the **SELECT** statement.

The corresponding expressions in the **SELECT** lists of the component queries of a compound query must match in number and data type.

1.7. Controlling the order of rows

By default, the output is sorted in ascending order on the first column. You can use the **ORDER BY** clause to change this. The **ORDER BY** clause can be used only once in a compound query. If used, the **ORDER BY** clause must be placed at the end of the query. The **ORDER BY** clause accepts the column name, an alias, or the positional notation.

Example:

```
SQL> COLUMN          a_dummy NOPRINT
SQL> SELECT          'sing' AS "My dream", 3 a_dummy
 2 FROM              dual
 3 UNION
 4 SELECT            'I'd like to teach', 1
 5 FROM              dual
 6 UNION
 7 SELECT            'the world to', 2
 8 FROM              dual
 9 ORDER BY          2;

My dream
-----
I'd like to teach
the world to
sing
```

→ The column **A_DUMMY** is useful to sort the result to display a correct sentence. The **NOPRINT** operator is used to not display the column with the sort criteria.

2. Oracle Server datetime functions

2.1. Datetime functions

2.1.1. TZ_OFFSET

The **TZ_OFFSET** function returns the time zone offset corresponding to the value entered. The value returned by this function is dependent on the date when the statement is executed.

For example, if the function returns -08:00, the value can be interpreted as the time zone from where the command was executed which is eight hours after UTC.

It's possible to enter a valid time zone name, a time zone offset from UTC, or the keyword **SESSIONTIMEZONE** or **DBTIMEZONE**.

Syntax:

TZ_OFFSET (*['time_zone_name'] '[+ | -] hh:mm' [SESSIONTIMEZONE] [DBTIMEZONE]*);

Example:

- The time zone 'US/Eastern' is four hours behind UTC.

```
SQL> SELECT TZ_OFFSET('US/Eastern') FROM DUAL;

TZ_OFFSET
-----
-04:00
```

- The time zone 'Canada/Yukon' is seven hours behind UTC.

```
SQL> SELECT TZ_OFFSET('Canada/Yukon') FROM DUAL;

TZ_OFFSET
-----
-07:00
```

You can query the V\$TIMEZONE-NAMES dynamic performance view, to list the valid time zone.

```
SQL> DESCRIBE v$tzone_names;

  Nom                                NULL ?    Type
-----
TZNAME                                VARCHAR2 (64)
TZABBREV                              VARCHAR2 (64)
```

2.1.2. CURRENT_DATE

The **CURRENT_DATE** function returns the current date in the session's time zone. Notice that the returned value is a date in the Gregorian calendar.

The following examples illustrate that **CURRENT_DATE** is sensitive to the session time zone. In the first statement, the session is altered to set the **TIME_ZONE** parameter to -5:0. The **TIME_ZONE** is a session parameter only, not an initialization parameter, which is set as follows:

TIME_ZONE = '[+ | -] hh:mm';

The format mask indicates the hours and minutes before or after UTC (Greenwich Mean Time) The **CURRENT_DATE** value changes when the **TIME_ZONE** parameter value is changed.

The **ALTER SESSION** command sets the date format of the session to 'DD-MON-YYY HH24: MI: SS'.

```
SQL> ALTER SESSION
  2  SET NLS_DATE_FORMAT = 'DD-MON-YYY HH24:MI:SS';

Session altered.

SQL> ALTER SESSION
  2  SET TIME_ZONE = '-5:0';

Session altered.

SQL> SELECT SESSIONTIMEZONE, CURRENT_DATE
  2  FROM DUAL;

SESSIONTIMEZONE          CURRENT_DATE
-----
-05:00                   31-AOU-004 04:29:20

SQL> ALTER SESSION SET TIME_ZONE = '-8:0';

Session altered.

SQL> SELECT SESSIONTIMEZONE, CURRENT_DATE
  2  FROM DUAL;

SESSIONTIMEZONE          CURRENT_DATE
-----
-08:00                   31-AOU-004 01:29:52
```

2.1.3.CURRENT_TIMESTAMP

This function returns the current date and time in the session time zone, as a value of the data type **TIMESTAMP WITH TIME ZONE**.

Syntax:

CURRENT_TIMESTAMP (*precision*);

Precision: Is an optional argument that specifies the fractional second precision of the time value returned. Default precision is 6.

The following examples illustrate that **CURRENT_TIMESTAMP** is sensitive to the session time zone.

First, the session is altered to set the **TIME_ZONE** parameter to -5:0.

Note that the **CURRENT_TIMESTAMP** value changes when the **TIME_ZONE** parameter is changed.

```
SQL> ALTER SESSION
      2 SET TIME_ZONE='-5:0';

Session altered.

SQL> SELECT SESSIONTIMEZONE, CURRENT_TIMESTAMP
      2 FROM DUAL;

SESSIONTIMEZONE          CURRENT_TIMESTAMP
-----
-05:00                  31/08/04 04:39:23,817000 -05:00

SQL> ALTER SESSION
      2 SET TIME_ZONE='-8:0';

Session altered.

SQL> SELECT SESSIONTIMEZONE, CURRENT_TIMESTAMP
      2 FROM DUAL;

SESSIONTIMEZONE          CURRENT_TIMESTAMP
-----
-08:00                  31/08/04 01:40:28,319000 -08:00
```

2.1.4.LOCALTIMESTAMP

The **LOCALTIMESTAMP** function returns the current date and time in the session time zone as a value of data type **TIMESTAMP**.

The difference between this function and the **CURRENT_TIMESTAMP** function is that it returns a **TIMESTAMP** value.

Syntax:

LOCAL_TIMESTAMP (*TIMESTAMP_precision*);

TIMESTAMP_precision: Is an optional argument that specifies the fractional second precision of the **TIMESTAMP** value returned by the function.

These examples illustrate the difference between **LOCALTIMESTAMP** and **CURRENT_TIMESTAMP**.

Note that the **LOCALTIMESTAMP** does not display the time zone value.

```
SQL> ALTER SESSION
2  SET TIME_ZONE='-5:0';

Session altered.

SQL> SELECT CURRENT_TIMESTAMP, LOCALTIMESTAMP
2  FROM DUAL;

CURRENT_TIMESTAMP                                LOCALTIMESTAMP
-----
31/08/04 04:47:23,217000 -05:00                31/08/04 04:47:23,217000

SQL> ALTER SESSION
2  SET TIME_ZONE='-8:0';

Session altered.

SQL> SELECT CURRENT_TIMESTAMP, LOCALTIMESTAMP
2  FROM DUAL;

CURRENT_TIMESTAMP                                LOCALTIMESTAMP
-----
31/08/04 01:48:00,304000 -08:00                31/08/04 01:48:00,304000
```

2.1.5.DBTIMEZONE and SESSIONTIMEZONE

The default database time zone is the same as the one of the operating system. This time zone is set by specifying the **SET TIME_ZONE** clause in the **CREATE DATABASE** statement.

The **DBTIMEZONE** function returns the value of the database time zone. The value returned is time zone offset, or a time zone region name, depending on how the user specified the database time zone value in the most recent **CREATE DATABASE** or **ALTER DATABASE** statement.

```
SQL> SELECT DBTIMEZONE FROM DUAL;

DBTIME
-----
+00:00
```

The **SESSIONTIMEZONE** function returns the value of the current session's time zone. The value returned in a time zone offset, or a time zone region name, depends on how the user specified the database time zone value in the most recent **CREATE DATABASE** or **ALTER DATABASE** statement.

Note that the database time zone is different from the current session's time zone.

```
SQL> SELECT          SESSIONTIMEZONE
2  FROM              DUAL;

SESSIONTIMEZONE
-----
-08:00
```

2.1.6.EXTRACT

The **EXTRACT** expression extracts and returns the value of a specified datetime field from a datetime or interval value expression.

The user can extract any of components mentioned in the following syntax.

Syntax:

```
SELECT EXTRACT ([YEAR] [MONTH] [DAY] [HOUR]
[MINUTE]
[SECOND] [TIMEZONE_HOUR]
[TIMEZONE_MINUTE]
[TIMEZONE_REGION] [TIMEZONE_ABBR]

FROM           [datetime_value_expression]
               [interval_value_expression]);
```

When extracting a **TIMEZONE_REGION** or **TIMEZONE_ABBR**, the value returned is a string containing the appropriate time zone name or abbreviation.

When extracting any other values the value returned is in the UTC format.

In the following statement, the **EXTRACT** function is used to extract the YEAR from **SYSDATE**.

```
SQL> SELECT      EXTRACT (YEAR FROM SYSDATE)
      FROM        DUAL;

EXTRACT (YEARFROMSYSDATE)
-----
                2004
```

In this example, the **EXTRACT** function is used to extract the MONTH from HIREDATE column of the EMP table, for all the employees.

```

SQL> SELECT      ename, hiredate,
2              EXTRACT (MONTH FROM hiredate)
3  FROM          emp;

```

ENAME	HIREDATE	EXTRACT (MONTHFROMHIREDATE)
SMITH	17-DEC-980 00:00:00	12
ALLEN	20-FEV-981 00:00:00	2
WARD	22-FEV-981 00:00:00	2
JONES	02-AVR-981 00:00:00	4
MARTIN	28-SEP-981 00:00:00	9
BLAKE	01-MAI-981 00:00:00	5
CLARK	09-JUN-981 00:00:00	6
SCOTT	09-DEC-982 00:00:00	12
KING	17-NOV-981 00:00:00	11
TURNER	08-SEP-981 00:00:00	9
ADAMS	12-JAN-983 00:00:00	1

ENAME	HIREDATE	EXTRACT (MONTHFROMHIREDATE)
JAMES	03-DEC-981 00:00:00	12
FORD	03-DEC-981 00:00:00	12
MILLER	23-JAN-982 00:00:00	1
GEROGES	26-AOU-004 11:24:58	8

15 row(s) selected.

2.2.Conversions

2.2.1.TIMESTAMP conversion using FROM_TZ

The **FROM_TZ** function is used to convert a **TIMESTAMP** value to a **TIMESTAMP WITH TIME ZONE** value.

Syntax:

FROM_TZ (**TIMESTAMP** *timestamp_value*, *time_zone_value*);

Time_zone_value: Is a character string in format 'TZH: TZM' or a character expression that returns a string **TZR** with optional **TZD** format.

TZR represents the time zone region in a datetime input strings.

The following example converts a **TIMESTAMP** value to a **TIMESTAMP WTH TIME ZONE VALUE**.

```

SQL> SELECT FROM_TZ (TIMESTAMP
2              '2000-03-28 08:00:00',
3              'Australia/North')
4  FROM      DUAL;

```

FROM_TZ (TIMESTAMP'2000-03-2808:00:00', 'AUSTRALIA/NORTH')
28/03/00 08:00:00,000000000 AUSTRALIA/NORTH

2.2.2.STRING to TIMESTAMP using TO_TIMESTAMP and TO_TIMESTAMP_TZ

The **TO_TIMESTAMP** function converts a string of CHAR, VARCHAR, NCHAR, or NVARCHAR2 data type to a value of TIMESTAMP data type.

Syntax:

TO_TIMESTAMP (*CHAR*, [*fmt*], [*'nlsparam'*]);

Fmt: Specifies the format of CHAR. If you omit it, the string must be in default format on the TIMESTAMP data type.

Nlsparam: Specifies the language in which the month and day names and abbreviations are returned. It can have this form: 'NLS_DATE_LANGUAGE= language'. If omitted, the function uses the default date language for your session.

```
SQL> SELECT      TO_TIMESTAMP('2000-12-01 11:00:00',
2                'YYYY-MM-DD HH:MI:SS') AS tmp_conversion
3  FROM          DUAL;

TMP_CONVERSION
-----
01/12/00 11:00:00,000000000
```

The **TO_TIMESTAMP_TZ** function converts a string of CHAR, VARCHAR, NCHAR, or NVARCHAR2 data type to a value of TIMESTAMP WITH TIME ZONE data type.

Syntax:

TO_TIMESTAMP_TZ (*CHAR*, [*fmt*], [*'nlsparam'*]);

Fmt: Specifies the format of CHAR. If you omit it, the string must be in default format on the TIMESTAMP WITH TIME ZONE data type.

Nlsparam: Specifies the language in which the month and day names and abbreviations are returned. It can have this form: 'NLS_DATE_LANGUAGE= language'. If omitted, the function uses the default date language for your session.

```
SQL> SELECT      TO_TIMESTAMP_TZ('1999-12-01 11:00:00 -8:00',
2                'YYYY-MM-DD HH:MI:SS TZh:TzM') AS tz_conversion
3  FROM          DUAL;

TZ_CONVERSION
-----
-
01/12/99 11:00:00,000000000 -08:00
```

2.2.3.Time interval conversion with TO_YMINTERVAL

The **TO_YMINTERVAL** function is used to convert a character string of CHAR, VARCHAR2, NCHAR, or NVARCHAR2 data type to an INTERVAL YEAR TO MONTH data type. This data type stores a period of time using the YEAR and MONTH datetime fields.

Syntax:

TO_YMINTERVAL (char);

Char: Is the character string to be converted.

The following statement calculates a date that is one year two months after the hire date for the employees working in the department 20 of the EMP table.

```
SQL> SELECT      hiredate,
2             hiredate + TO_YMINTERVAL('01-02') AS
3             hire_date_YMIN
4 FROM          emp
5 WHERE         deptno=20;
```

HIREDATE		HIRE_DATE_YMIN
-----		-----
17-DEC-1980 00:00:00		17-FEV-1982 00:00:00
02-AVR-1981 00:00:00		02-JUN-1982 00:00:00
09-DEC-1982 00:00:00		09-FEV-1984 00:00:00
12-JAN-1983 00:00:00		12-MAR-1984 00:00:00
03-DEC-1981 00:00:00		03-FEV-1983 00:00:00

3. GROUP BY with ROLLUP and CUBE operators

3.1. Reviews of GROUP BY and HAVING

3.1.1. Review of group functions

See. SQLP course “Module 2: Techniques of data retrieval”.

3.1.2. Review of the GROUP BY clause

See. SQLP course “Module 2: Techniques of data retrieval”.

3.1.3. Review of the HAVING clause

See. SQLP course “Module 2: Techniques of data retrieval”.

3.2. ROLLUP

3.2.1. The ROLLUP operator

The **ROLLUP** operator is an extension of the **GROUP BY** clause that can produce cumulative aggregates such as sub-totals.

```
SELECT      column, group_function
FROM        table
[WHERE      condition]
[GROUP BY   [ROLLUP] (group_by_expression)];
```

The **ROLLUP** operator creates groups by moving in one direction, from right to left, along the list of columns specified in the **GROUP BY** clause. It then applies the aggregate function to these groups.

3.2.2. Example

The **ROLLUP** operator creates subtotals that roll up from the most detailed level to a grand total, following the grouping list specified in the **GROUP BY** clause. First it computes the standard aggregate values for the groups specified in the **GROUP BY** clause. Then it creates progressively higher-level subtotals, moving from right to left through the list of grouping columns.

Given two expressions in the **ROLLUP** operator of the **GROUP BY** clause, the operation results in $n+1=2+1=3$ groupings.

Rows based on the values of the first n expressions are called rows or regular rows and the others are called super aggregate rows.

Example:

```
SQL> SELECT      deptno,job, SUM(sal)
  2 FROM          emp
  3 GROUP BY      ROLLUP (deptno,job);
```

DEPTNO	JOB	SUM(SAL)
10	CLERK	1300
10	MANAGER	2450
10	PRESIDENT	5000
10		8750
20	ANALYST	6000
20	CLERK	1900
20	MANAGER	2975
20		10875
30	CLERK	950
30	MANAGER	2850
30	SALESMAN	5600
30		9400
		29025

13 row(s) selected.

-> This statement displays the sum of all salaries for each job in each department, as well as sum of all salaries for each department and for all of them.

3.3.CUBE

3.3.1.The CUBE operator

The **CUBE** operator is an additional switch in the **GROUP BY** clause in the **SELECT** statement. The **CUBE** operator can be applied to all aggregate functions.

```
SELECT      column, group_function
FROM        table
[WHERE      condition]
[GROUP BY   [CUBE] (group_by_expression)];
```

While **ROLLUP** produces only a fraction of possible subtotal combinations, **CUBE** produces subtotals for all possible combinations of groupings specified in the **GROUP BY** clause, and a grand total.

All the returned values are calculated from the specified aggregate group in the **SELECT** list.

3.3.2.Example

CUBE operator returns the same result as **ROLLUP** operator, but its group function can be applied to subgroup also.

The number of additional groups in the result is determined by the number of columns in the **GROUP BY** clause, because each columns combination is used to produce supersets. Then if there is n column(s) or expressions in the **GROUP BY** clause, there will be 2^n possible combinations supersets.

Example:

```
SQL> SELECT      deptno, job, SUM(sal)
2  FROM          emp
3  GROUP BY CUBE (deptno, job);
```

DEPTNO	JOB	SUM(SAL)
10	CLERK	1300
10	MANAGER	2450
10	PRESIDENT	5000
10		8750
20	ANALYST	6000
20	CLERK	1900
20	MANAGER	2975
20		10875
30	CLERK	950
30	MANAGER	2850
30	SALESMAN	5600
30		9400
	ANALYST	6000
	CLERK	4150
	MANAGER	8275
	PRESIDENT	5000
	SALESMAN	5600
		29025

→ This statement returns the same result as the **ROLLUP** operator. The sum of the salary has been added.

3.4. Analytical functions

3.4.1. Describe analytical functions

To ease advanced SQL programming, analytical functions were introduced since Oracle8i Release 2, to simplify calculations such as average variables and classifications.

Groups defined by the **GROUP BY** clause of a **SELECT** are called partitions. The result of a statement can be composed by a partition with all the lines, some big partitions, or many small partitions with few lines each. Analytical functions are applied to each line of each partition.

3.4.2. RANK function

The **RANK** function creates a classification starting with 1.

```
SELECT      column, group_function(argument),
            RANK() OVER([PARTITION BY column] ORDER BY
                        group_function(argument [DESC])
FROM table
GROUP BY    column
```

Example:

```
SQL> SELECT      deptno, job, SUM(sal),
2              RANK() OVER(PARTITION BY deptno
3              ORDER BY SUM(sal) DESC)
4              AS rank_of_job_per_dep,
5              RANK() OVER(ORDER BY SUM(sal) DESC)
6              AS rank_of_sumsal
7 FROM          emp
8 GROUP BY      deptno, job
9 ORDER BY      deptno;
```

DEPTNO	JOB	SUM(SAL)	RANK_OF_JOB_PER_DEP	RANK_OF_SUMSAL
10	ANALYST	5000	1	3
10	MANAGER	2450	2	6
10	CLERK	1300	3	8
20	ANALYST	6000	1	1
20	MANAGER	2975	2	4
20	CLERK	1900	3	7
30	SALESMAN	5600	1	2
30	MANAGER	2850	2	5
30	CLERK	950	3	9

9 row(s) selected.

-> This statement displays the sum of salaries classification by department and for each one of these.

The classification is applied on all the columns specified in the **ORDER BY**, if none is specified the classification will be on all columns. **RANK** gives a rank from 1 to the smallest value, except when **DESC** is used.

3.4.3.CUME_DIST function

The **CUME_DIST** function calculates the relative position of a value compared to the other values of the partition.

```
SELECT      column, group_function(argument),
            CUME_DIST() OVER([PARTITION BY column]
            ORDER BY group_function(argument) [DESC])
FROM        table
GROUP BY    column
```

The **CUME_DIST** function determines the part of the lines of the partition which are lower or equal to the current value. The result is a decimal value between zero and one. By default, the order is ascending, i.e. the smallest value of the partition corresponds to the smallest **CUME_DIST**.

Example:

```
SQL> SELECT      deptno, job, SUM(sal),
2              CUME_DIST() OVER(PARTITION BY deptno
3              ORDER BY SUM(sal) DESC)
4              AS cume_dist_per_dep
5 FROM          emp
6 GROUP BY      deptno, job
7 ORDER BY      deptno, SUM(sal);
```

DEPTNO	JOB	SUM(SAL)	CUME_DIST_PER_DEP
10	CLERK	1300	1
10	MANAGER	2450	,666666667
10	PRESIDENT	5000	,333333333
20	CLERK	1900	1
20	MANAGER	2975	,666666667
20	ANALYST	6000	,333333333
30	CLERK	950	1
30	MANAGER	2850	,666666667
30	SALESMAN	5600	,333333333

9 row(s) selected.

-> This statement displays **CUME_DIST** of the sum of the salaries for each job in each department.

3.5.GROUPING

3.5.1.The GROUPING function

When using **ROLLUP** and **CUBE** operator, blanks appear in the result. The **GROUPING** function exists to differentiate these blanks from NULL values. It can determine the sub-total level, i.e. the groups from which sub-totals are calculated.

```
SELECT      column, group_function, GROUPING(expr)
FROM        table
[WHERE      condition]
[GROUP BY   [ROLLUP][CUBE] (group_by_expression)];
```

The **GROUPING** function can take only one argument. This one must be the same as one of the expressions in the **GROUP BY** clause.

3.5.2.Example

The **GROUPING** function acts as a Boolean function.

A value of 0 is returned if:

- The expression has been used to calculate the aggregate value.
- The NULL value in the expression column is a stored NULL value.

A value of 1 is returned if:

- The expression has not been used to calculate the aggregate value.
- The NULL value in the expression column is created by **ROLLUP/CUBE** as a result of grouping.

Example:

```
SQL> SELECT      deptno, job, SUM(sal), GROUPING(deptno),
2              GROUPING(job)
3 FROM          emp
4 GROUP BY ROLLUP (deptno, job);
```

DEPTNO	JOB	SUM(SAL)	GROUPING(DEPTNO)	GROUPING(JOB)
10	CLERK	1300	0	0
10	MANAGER	2450	0	0
10	PRESIDENT	5000	0	0
10		8750	0	1
20	ANALYST	6000	0	0
20	CLERK	1900	0	0
20	MANAGER	2975	0	0
20		10875	0	1
30	CLERK	950	0	0
30	MANAGER	2850	0	0
30	SALESMAN	5600	0	0
30		9400	0	1
		29025	1	1

13 row(s) selected.

-> This statement displays two **GROUPING** columns that indicate if deptno and job have been calculated with the **ROLLUP** operator. When the value is 1, it shows that the column has been used and thus it is not a NULL value.

When **GROUPING** function is used with **CUBE** operator, the column **GROUPING(deptno)** will have a “1” and the column **GROUPING(job)** a “0”, because deptno is not used to calculate the sum of the sub-group.

3.5.3.GROUPING SETS

You can use the **GROUPING SETS** function to define multiple groupings in the same query. **GROUPING SETS** are a further extension of the **GROUP BY** clause that let you specify multiple groupings of data.

A single **SELECT** statement can now be written with the **GROUPING SETS** function to specify various groupings, rather than multiple **SELECT** statements combined by **UNION ALL** operators.

3.5.4.Using GROUPING SETS

```
SQL> SELECT      deptno, job, mgr, AVG (sal)
2  FROM          emp
3  GROUP BY      GROUPING SETS
4      ((deptno, job), (job, mgr));
```

DEPTNO	JOB	MGR	AVG (SAL)
10	CLERK		1300
10	MANAGER		2450
10	PRESIDENT		5000
20	ANALYST		3000
20	CLERK		950
20	MANAGER		2975
30	CLERK		950
30	MANAGER		2850
30	SALESMAN		1400
60	ANALYST	7566	3000

DEPTNO	JOB	MGR	AVG (SAL)
	CLERK	7698	950
	CLERK	7782	1300
	CLERK	7788	1100
	CLERK	7902	800
	MANAGER	7839	2758,33333
	PRESIDENT		5000
	SALESMAN	7698	1400

19 row(s) selected.

-> This query computes aggregates over two groupings.

The table is divided into two groups: deptno, job and job, mgr.

The average salary of each of these groups is calculated. The result set displays average salary for each of the two groups.

3.6.Composite columns

3.6.1.The composite columns

A composite column is a collection of columns treated as a unit during the computation of groupings. To use composite column, you have to specify the columns in parentheses as in the following example:

ROLLUP (column A, (column B, column C))

Here, (column B, column C) are treated as a unit and **ROLLUP** will not be applied across (column B, column C).

Composite columns are usually used in **ROLLUP**, **CUBE**, and **GROUPING SETS**.

GROUPIN SETS statements	Equivalent GROUP BY statements
GROUP BY GROUPING SETS (a, b, c)	GROUP BY a UNION ALL GROUP BY b UNION ALL GROUP BY c
GROUP BY GROUPING SETS (a, b, (b, c)) (The GROUPIN SET has a composite column)	GROUP BY a UNION ALL GROUP BY b UNION ALL GROUP BY b, c
GROUP BY GROUPING SETS ((a, b, c))	GROUP BY a, b, c
GROUP BY GROUPING SETS (a, (b, c))	GROUP BY a UNION ALL GROUP BY b UNION ALL GROUP BY (b, c)

3.6.2.Example

```
SQL> SELECT deptno, job, mgr, SUM(sal)
2 FROM emp
3 GROUP BY ROLLUP(deptno, (job, mgr));
```

DEPTNO	JOB	MGR	SUM(SAL)
10	CLERK	7782	1300
10	MANAGER	7839	2450
10	PRESIDENT		5000
10			8750
20	ANALYST	7566	6000
20	CLERK	7788	1100
20	CLERK	7902	800
20	MANAGER	7839	2975
20			10875
30	CLERK	7698	950
30	MANAGER	7839	2850
30	SALESMAN	7698	5600
30			9400
60			
60			
			29025

16 row(s) selected.

-> This statement computes the following groupings:

- (deptno, job, mgr)
- (deptno)
- ()

And displays:

- Total salary for every department.
- Total salary for every department , job, and manager
- Grand total.

3.7.Concatenated groupings

3.7.1.The concatenated groupings

Concatenated groupings offer a concise way to generate useful combinations of groupings. They are specified by listing multiple **GROUPING SETS**, **CUBE**, and **ROLLUP**, and separating them with commas.

GROUP BY GROUPING SETS (a, b), GROUPING SETS (c, d)

This preceding statement defines the following groupings:

(a, c), (a, d), (b, c), (b, d).

3.7.2.Example

```
SQL> SELECT      deptno, job, mgr, AVG (sal)
2  FROM          emp
3  GROUP BY      deptno,
4                ROLLUP (job) ,
5                CUBE (mgr) ;
```

DEPTNO	JOB	MGR	AVG (SAL)
10	CLERK	7782	1300
10	MANAGER	7839	2450
10	PRESIDENT		5000
20	ANALYST	7566	3000
20	CLERK	7788	1100
20	MANAGER	7839	2975
20	CLERK	7902	800
30	CLERK	7698	950
30	SALESMAN	7698	1400
30	MANAGER	7839	2850
60			
	...		

35 row(s) selected(s).

-> This statement results in the following groupings:

- (deptno, mgr, job)
- (deptno, mgr)
- (deptno, job)
- (deptno).

The total salary for each of these groups is calculated.

This example displays the following:

- Total salary for every department, manager, and job.
- Total salary for every department and manager.
- Total salary for every department and job.
- Total salary for every department.

4. Advanced subqueries

4.1. Subqueries

4.1.1. What is a subquery

See. SQLP course “Module 2: Techniques of data retrieval”

4.1.2. Multiple-column subqueries

See. SQLP course “Module 2: Techniques of data retrieval”

4.1.3. Using subqueries

See. SQLP course “Module 2: Techniques of data retrieval”

4.2. Column comparisons

4.2.1. Pairwise comparisons

This SQL statement contains a multiple-column subquery because the subquery returns more than one column.

It compares that value in the MGR column and the DEPTNO column of each row in the EMP table with the values in the MGR column and the DEPTNO column for the employees with EMPNO 7369 or 7499.

First, the subquery retrieves the MGR and DEPTNO values for the employees with EMPNO 7369 or 7499.

Then the values are compared with the MGR column and the DEPTNO column for each row in the EMP table.

If the values match, the row is displayed; the employee number 7369 and 7499 will not be displayed.

SQL>	SELECT	empno, mgr, deptno
2	FROM	emp
3	WHERE	(mgr, deptno) IN
4		(SELECT mgr, deptno
5		FROM emp
6		WHERE empno IN (7369,7499))
7	AND	empno NOT IN (7369,7499);
	EMPNO	MGR DEPTNO
	-----	-----
	7521	7698 30
	7844	7698 30
	7900	7698 30
	7654	7698 30

4.2.2. Nonpairwise comparisons

The following example shows a nonpairwise comparison.

It displays the EMPNO, MGR, and DEPTNO of any employee whose manager ID matches any of the manager IDs of employees whose employee IDs are either 7639 or 7499 and DEPTNO matches any of the department IDs of employees whose employee IDs are either 7369 or 7499.

First, the subquery retrieves the MGR values for the employees with the EMPNO 7369 or 7499. The second subquery retrieves the DEPTNO values for the same employees.

The retrieved values of the MGR and DEPTNO columns are compared with the MGR and DEPTNO columns for each row in the EMP table.

If the MGR column of the row in the EMP table matches with any of the values of the MGR retrieved by the inner subquery and if the DEPTNO column of the row in the EMP table matches with any of the values of the DEPTNO retrieved by the second subquery, the record is displayed.

SQL>	SELECT	empno, mgr, deptno
2	FROM	emp
3	WHERE	mgr IN
4		(SELECT mgr
5		FROM emp
6		WHERE empno IN (7369,7499))
7	AND	deptno IN
8		(SELECT deptno
9		FROM emp
10		WHERE empno IN (7369,7499))
11	AND	empno NOT IN (7369,7499);

EMPNO	MGR	DEPTNO
7521	7698	30
7844	7698	30
7900	7698	30
7654	7698	30

4.3. Scalar subqueries

4.3.1. The scalar subqueries

We qualify as a scalar subquery, all subqueries that return exactly one column value from one row. Multiple-column subqueries written to compare two or more columns, using a compound **WHERE** clause and logical operators do not qualify as scalar subqueries.

The value of the scalar subquery expression is the value of the select list item of the subquery. If the subquery returns zero rows, the value of the scalar subquery expression is NULL. If it returns more than one row, the Oracle Server returns an error.

You can use scalar subqueries in:

- Condition and expression part of **DECODE** and **CASE**.
- All clauses of **SELECT** except **GROUP BY**.
- In the left-hand side of the operator in the **SET** clause and **WHERE** clause of **UPDATE**.

Scalar subqueries are not valid in the following places:

- As default values for columns and hash expressions for clusters.
- In the **RETURNING** clause of DML statements.
- As the basis of a function-based index.

- In **GROUP BY** clauses, **CHECK** constraints, **WHEN** condition.
- **HAVING** clauses.
- In **START WITH** and **CONNECT BY** clauses.
- In statements that are unrelated to queries, such as **CREATE PROFILE**.

4.3.2.Example

In the following statement, the inner query returns the value 20, which is the department ID of the department located in DALLAS.

The **CASE** expression in the outer query uses the result of the inner query to display the employee ID, name, and a value of Canada or USA, depending on whether the department ID of the record retrieved by the outer query is 20 or not.

```
SQL> SELECT      empno,ename,
2              (CASE
3              WHEN deptno=
4                  (SELECT deptno FROM dept
5                   WHERE LOC='DALLAS')
6              THEN ' USA' ELSE 'Canada ' END) location
7 FROM          emp;
```

EMPNO	ENAME	LOCATI
7369	SMITH	Canada
7499	ALLEN	USA
7521	WARD	USA
7566	JONES	Canada
7654	MARTIN	USA
7698	BLAKE	USA
7782	CLARK	USA
7788	SCOTT	Canada
7839	KING	USA
7844	TURNER	USA
7876	ADAMS	Canada

EMPNO	ENAME	LOCATI
7900	JAMES	USA
7902	FORD	Canada
7934	MILLER	USA
7777	GEROGES	USA

15 row(s) selected.

4.4.Correlated subqueries

4.4.1.The correlated subqueries

A correlated subquery is a nested subquery that is evaluated once for each row processed by the main query and which on execution uses a value from a column in the outer query. The correlation can be obtained by using an element from the main query in the subquery.

Stages of execution of a correlated query:

- Get a candidate row (fetched by the outer query).
- Execute the inner query using the value of the candidate row.
- Use the values resulting from the inner query to qualify or disqualify the candidate.
- Repeat until no candidate row remains.

Although this discussion focuses on correlated subqueries in **SELECT** statements, it also applies to correlated **UPDATE** and **DELETE** statements.

4.4.2.Example

A correlated subquery is the mean to read every row in a table and compare values in each row with related data.

```
SELECT      outer1, outer2, ...
FROM        table1 alias1
WHERE       outer1 operator
            (SELECT      inner1
             FROM         table2 alias2
             WHERE        alias1.expr1 =
                        alias2.expr2);
```

It is used whenever a subquery must return a different result or set of results for each candidate row considered by the main query. In other words, you use a correlated subquery to answer a multipart question whose answer depends on the value in each row processed by the parent statement. The Oracle server performs a correlated subquery when the subquery references a column from a table in the parent query.

Example:

```
SQL> SELECT      empno, sal, deptno
2  FROM          emp outer
3  WHERE         sal > (SELECT      AVG(sal)
4                      FROM          emp inner
5                      WHERE         outer.deptno = inner.deptno);
```

EMPNO	SAL	DEPTNO
7499	1600	30
7566	2975	20
7698	2850	30
7788	3000	20
7839	5000	10
7902	3000	20

6 row(s) selected.

→ This statement displays the entire list of employees that have a salary superior to the average salary of their department.

When the main and the correlated query refer to the same table, you have to use a table alias to avoid false results.

4.5.EXITS and NOT EXISTS operators

4.5.1.The EXISTS operator

With nesting SELECT statements, all logical operators are valid. In addition, you can use the EXISTS operator. This operator is frequently used with correlated subqueries to test whether a value is there.

If the subquery returns a line:

- The search doesn't continue in the subquery.
- The condition is labelled TRUE.

If the subquery doesn't return a value for a line:

- The condition is labelled FALSE.
- The search continues with the following line in the subquery.

4.5.2.Using the EXISTS operator

Example:

```
SQL> SELECT      empno, ename, job, deptno
  2 FROM          emp outer
  3 WHERE         EXISTS (SELECT      'X'
  4                      FROM        emp inner
  5                      WHERE inner.mgr = outer.empno);
```

EMPNO	ENAME	JOB	DEPTNO
7566	JONES	MANAGER	20
7698	BLAKE	MANAGER	30
7782	CLARK	MANAGER	10
7788	SCOTT	ANALYST	20
7839	KING	PRESIDENT	10
7902	FORD	ANALYST	20

6 row(s) selected

→ This statement displays the entire list of employees that have at least one person under their orders.

Note that the inner SELECT query does not need to return a specific value, so a literal can be selected. From a performance standpoint, it is faster to select a constant than a column.

4.5.3. The NOT EXISTS operator

Example:

```
SQL> SELECT      deptno, dname
  2 FROM          dept d
  3 WHERE         NOT EXISTS (SELECT      'X'
  4                      FROM        emp e
  5                      WHERE        d.deptno = e.deptno);
```

DEPTNO	DNAME
40 OPERATIONS	

→ This statement displays all the departments in which there is no employee.

A NOT IN construct can be used as an alternative for a NOT EXISTS operator. However, use it with caution. NOT IN evaluates to FALSE if any member of the set is a NULL value. Therefore, your query will not return any rows.

Example:

```
SQL> SELECT      o.ename
  2 FROM emp o
  3 WHERE      NOT EXISTS (SELECT  'X'
  4                      FROM      emp i
  5                      WHERE      i.mgr = o.empno);

ENAME
-----
SMITH
ALLEN
WARD
MARTIN
TURNER
ADAMS
JAMES
MILLER

8 row(s) selected.
```

→ This statement displays all the employees that have anybody to give orders to.

The statement from the example can also be written with a NOT IN clause but the result will differ because NULL values will be used.

Example:

```
SQL> SELECT      o.ename
  2 FROM emp o
  3 WHERE      o.empno NOT IN (SELECT      i.mgr
  4                      FROM      emp i);

No lines selected.
```

→ This statement doesn't return anything because the subquery returns a NULL value and all comparison with a NULL value returns a NULL value; the result is totally different.

When the subquery returns NULL values, it is better not to use NOT IN as an alternative to NOT EXISTS.

4.6. Correlated UPDATE and DELETE

4.6.1. The correlated UPDATE

In the case of the UPDATE statement, you can use a correlated subquery to update rows in one table based on rows from another table.

```
UPDATE      table1 alias1
SET  column = (SELECT expression
                FROM table2 alias2
                WHERE alias1.column = alias2.column);
```

Example:

```
SQL> ALTER TABLE emp
2  ADD      (dname VARCHAR2(14));

SQL> UPDATE      emp e
2  SET      dname = (SELECT      dname
3                      FROM      dept d
4                      WHERE      e.deptno = d.deptno);

14 row(s) updated.
```

4.6.2.The correlated DELETE

In the case of a DELETE statement, you can use a correlated subquery to delete only those rows that also exist in another table.

```
DELETE FROM      table1 alias1
WHERE      column operator
            (SELECT expression
              FROM table2 alias2
              WHERE alias1.column = alias2.column);
```

Example:

```
SQL> DELETE FROM emp E
2  WHERE      empno =
3              (SELECT      empno
4                      FROM      emp_history EH
5                      WHERE      E.empno = EH.empno);
```

→ This statement deletes from EMP table the lines that are already on the EMP_HISTORY table.

4.7.WITH clause

4.7.1.The WITH clause

Using the **WITH** clause, you can define a query block before using it in a query.

With this clause, you can reuse the same query in a **SELECT** statement when it is used more than once within a complex query.

The Oracle Server retrieves the results of a query block and stores it in the user's temporary tablespace.

This improves performance.

4.7.2.Example

The following example represents a situation in which it's possible to improve the performance and write SQL more simply by using the **WITH** clause.

The query creates the query names DEPT_COST and AVG_COST and then uses them in the body of the main query.

Internally, the **WITH** clause is resolved either as an inline view or a temporary table.

```
SQL> WITH
  2 dept_cost AS (
  3   SELECT    d.dname, SUM(e.sal) AS dept_total
  4   FROM      emp e, dept d
  5   WHERE     e.deptno=d.deptno
  6   GROUP BY  d.dname),
  7 avg_cost AS (
  8   SELECT    SUM(dept_total)/COUNT(*) AS dept_avg
  9   FROM      dept_cost)
10 SELECT
11 FROM      dept_cost
12 WHERE     dept_total >
13           (SELECT dept_avg
14            FROM   avg_cost)
15 ORDER BY  dname;
```

DNAME	DEPT_TOTAL
ACCOUNTING	8750
RESEARCH	10875
SALES	9400

5. Hierarchical retrieval

5.1. Hierarchical queries overview

5.1.1. When is a hierarchical query possible?

Hierarchical queries enable you to retrieve data based on a natural hierarchical relationship between rows in a table.

A relational database does not store records in a hierarchical way. However, where a hierarchical relationship exists between the rows of a single table, a process called tree walking enables the hierarchy to be constructed. A hierarchical query is a method of reporting, in order, the branches of a tree.

5.1.2. Natural structure tree

A tree structure is composed by a parent node that is divided into child branch, which can be divided into branches and so on. For example, the EMP table has a tree structure representing the management reporting line. The hierarchy can be created by looking at the relationship between equivalent values in the EMPNO and MGR columns. This relationship has been exploited by joining the table to itself. An employee's MGR number is the employee number of his or her manager.

The parent-child relationship of a tree structure enables you to control:

- The direction in which the hierarchy is walked
- The starting point inside the hierarchy

5.1.3. Hierarchical queries

Hierarchical queries are queries that contain CONNECT BY and START WITH clauses.

```
SELECT          [LEVEL], column, expr...
FROM            table
[WHERE          condition(s)]
[START WITH     condition(s)]
[CONNECT BY PRIOR condition(s)];
```

LEVEL is a pseudo column, which returns 1 for the root node, 2 for a child of the root and so on. **LEVEL** counts how far down a hierarchical tree you have travelled.

START WITH specifies the root rows of the hierarchy. This clause is required for a true hierarchical query.

CONNECT BY PRIOR specifies the columns in which the relationship between parent and child rows exists. This clause is required for a hierarchical query.

The **SELECT** statement cannot contain a join or a query from a view that contains a join.

5.2. Walking the tree

5.2.1. Start point

The row or rows to be used as the root of the tree are determined by the **START WITH** clause. The **START WITH** clause can be used in conjunction with any valid condition and can contain a sub query.

If the **START WITH** clause is omitted, the tree walk is started with all the rows in the table as root rows. If a **WHERE** clause is used, the walk is started with all the rows that satisfy the **WHERE** condition. This no longer reflects a true hierarchy.

The **START WITH** and **CONNECT BY PRIOR** are not a part of the standard ANSI SQL.

5.2.2.Direction

The hierarchical direction of the query, whether it is from parent to child or from child to parent, is determined by the **CONNECT BY PRIOR** column placement. The **PRIOR** operator refers to the parent row. To find the children of a parent row, the Oracle server evaluates the **PRIOR** expression for the parent row and the other expressions for each row in the table. Rows for which the condition is true are the children of the parent. The Oracle server always selects children by evaluating the **CONNECT BY** condition with respect to a current parent row.

For example, to walk from the top down using the EMP table you have to define a hierarchical relationship in which the EMPNO value of the parent row is equal to the MGR value of the child row.

```
... CONNECT BY PRIOR empno=mgr
```

The **CONNECT BY** clause cannot contain subqueries.

5.2.3.Example

Example:

```
SQL> SELECT      empno, ename, job, mgr
2  FROM          emp
3  START WITH    empno = 7698
4  CONNECT BY    PRIOR mgr = empno;
```

EMPNO	ENAME	JOB	MGR
7698	BLAKE	MANAGER	7839
7839	KING	PRESIDENT	

-> This statement displays a tree walk from top to bottom of the EMP table, starting with the employee 7698.

Expressions in the conditions can represent more than a single column. In that case, **PRIOR** operator is only valid for the first one.

Example:

```
SQL> SELECT      empno, ename, job, mgr
2  FROM          emp
3  START WITH    empno = 7698
4  CONNECT BY    PRIOR empno = mgr AND sal > NVL(comm,0);
```

EMPNO	ENAME	JOB	MGR
7698	BLAKE	MANAGER	7839
7499	ALLEN	SALESMAN	7698
7521	WARD	SALESMAN	7698
7844	TURNER	SALESMAN	7698
7900	JAMES	CLERK	7698

-> This statement returns the lines which manager number equal to employee number of the parent line, and which salary is superior from the commission.

5.3. Organizing data

5.3.1.Ranking rows with the LEVEL pseudo column

You can explicitly show the rank or level of a row in the hierarchy by using the **LEVEL** pseudo column. This will make your report more readable.

The point where one or more branches split away from a larger branch are called nodes, and the very end of a branch is called a leaf, or leaf node. The value of the **LEVEL** pseudo column depends on the

node on which the line is. The value of the **LEVEL** pseudo column will be 1 for the parent node, then 2 for a child node and so on.

The first node of a tree is called the root node, whereas others are called child node. A parent node is a node that has at least one child. A leaf node is a node without any child.

5.3.2. Formatting hierarchical report using LEVEL and LPAD

It is possible to graphically reflect the hierarchy: use the LPAD function in conjunction with the pseudo column **LEVEL** to display a hierarchical report as an indented tree (cf. SQLP course “Module 1: Basic **SELECT** statements”).

Example:

```
SQL> COLUMN org_chart FORMAT A15
SQL> SET PAGESIZE 20
SQL> SELECT          LPAD(' ', 3 * LEVEL-3)||ename AS org_chart,
2                  LEVEL, empno, mgr, deptno
3 FROM              emp
4 START WITH        mgr IS NULL
5 CONNECT BY PRIOR empno = mgr;
SQL> CLEAR COLUMN
```

ORG_CHART	LEVEL	EMPNO	MGR	DEPTNO
KING	1	7839		10
JONES	2	7566	7839	20
SCOTT	3	7788	7566	20
ADAMS	4	7876	7788	20
FORD	3	7902	7566	20
SMITH	4	7369	7902	20
BLAKE	2	7698	7839	30
ALLEN	3	7499	7698	30
WARD	3	7521	7698	30
MARTIN	3	7654	7698	30
TURNER	3	7844	7698	30
JAMES	3	7900	7698	30
CLARK	2	7782	7839	10
MILLER	3	7934	7782	10

14 row(s) selected.

-> This statement adds spaces in front of the employee names according to their level in the hierarchy.

5.3.3. Pruning branches

In a hierarchical query, it is possible to delete a branch or a node of a tree.

To prune a branch, you must use the **WHERE** clause. Thus the concerned branch is not shown but child branches are.

Example:

```
SQL> SELECT      deptno, empno, ename, job, sal
2  FROM          emp
3  WHERE         ename != 'SCOTT'
4  START WITH    mgr IS NULL
5  CONNECT BY    PRIOR empno = mgr;
```

DEPTNO	EMPNO	ENAME	JOB	SAL
10	7839	KING	PRESIDENT	5000
20	7566	JONES	MANAGER	2975
20	7876	ADAMS	CLERK	1100
20	7902	FORD	ANALYST	3000
20	7369	SMITH	CLERK	800
30	7698	BLAKE	MANAGER	2850
30	7499	ALLEN	SALESMAN	1600
30	7521	WARD	SALESMAN	1250
30	7654	MARTIN	SALESMAN	1250
30	7844	TURNER	SALESMAN	1500
30	7900	JAMES	CLERK	950
10	7782	CLARK	MANAGER	2450
10	7934	MILLER	CLERK	1300

-> This statement walks the tree, from top to bottom, from the root node without showing SCOTT, but including child branch (ADAMS).

5.3.4.Ordering Data

It is recommended that you do not use the **ORDER BY** clause when creating hierarchical query reports because the implicit natural ordering of the tree may be destroyed.

5.3.5.ROW_NUMBER() function

Default, NULL values are displayed first, when the value are sorted in a decreasing order. Since Oracle 8i Release 2, it is possible to force NULL values to be displayed last, by adding **NULLS LAST** in the **ORDER BY** clause.

The **ROW_NUMBER** function gives to each line of the partition, a number based on the sequence defined by the **ORDER BY** clause.

Example:

```
SQL> SELECT      empno, job, comm,
2              ROW_NUMBER() OVER(ORDER BY comm DESC NULLS LAST)
3              AS rnum
4 FROM          emp;
```

EMPNO	JOB	COMM	RNUM
7654	SALESMAN	1400	1
7521	SALESMAN	500	2
7499	SALESMAN	300	3
7844	SALESMAN	0	4
7369	CLERK		5
7566	MANAGER		6
7900	CLERK		7
7934	CLERK		8
7902	ANALYST		9
7876	CLERK		10
7698	MANAGER		11
7782	MANAGER		12
7788	ANALYST		13
7839	PRESIDENT		14

14 row(s) selected.

-> This statement numbers the lines according to the value of the commission and displays NULL values last

6.DML and DDL statements

6.1.Multitable INSERT statement

6.1.1.Types of multitable INSERT statement

You can insert a row into multiple tables as part of a single DML statement using the **INSERT...SELECT** statement.

Here are the different types of multitable **INSERT** statements:

- Unconditional **INSERT ALL**.
- Conditional **INSERT ALL**.
- Conditional **FIRST INSERT**.
- Pivoting **INSERT**.

Syntax:

```
INSERT [ALL] [Conditional_insert_clause]
[insert_into_clause Values_clause]      (Subquery)
```

Conditional_insert_clause:

```
[ALL] [FIRST]
[WHEN condition THEN] [insert_into_clause_values_clause]
[ELSE]                [Insert_into_clause Values_clause]
```

Syntax	Description
Unconditional INSERT:ALL into_clause	Specify ALL followed by multiple <i>insert_into_clauses</i> to perform an unconditional multitable insert. The Oracle Server executes each <i>insert_into_clause</i> once for each row returned by the subquery.
Conditional INSERT: conditional_insert_clause	Specify the <i>conditional_insert_clause</i> to perform a conditional multitable insert. The Oracle Server filters each <i>insert_into_clause</i> through the corresponding WHEN condition, which determines that <i>insert_into_clause</i> is executed. A single multitable insert statement can contain up to 127 WHEN clauses.
Conditional FIRST: INSERT	If you specify FIRST , the Oracle Server evaluates each WHEN clause in the order in which it appears in the statement. If the first WHEN clause is evaluated to true, the Oracle Server executes the corresponding INTO clause and skips subsequent WHEN clauses for the given row.

Conditional INSERT: ELSE clause

For a given row, if no **WHEN** clauses is evaluated to true, and if you have specified an **ELSE** clause, the Oracle Server executes the **INTO** clause list associated with the **ELSE** clause, but if you didn't specify an **ELSE** clause, the Oracle Server takes no action for that row.

6.1.2.Unconditional INSERT ALL

The following statement inserts rows into the SAL_HISTORY and the MGR_HISTORY table. The **SELECT** statement retrieves the details of employee ID, hire date, salary, and manager ID of those employees whose employee ID is greater than 7400 from the EMP table.

The details of the employee ID, hire date, and salary are inserted into the SAL_HISTORY table.

The details of employee ID, manager ID, and salary are inserted into the MGR_HISTORY table.

The **INSERT** statement is referred to as an unconditional **INSERT**, as no further restriction is applied to the rows that are retrieved by the **SELECT** statement. All the rows retrieved by the **SELECT** statement are inserted into the two tables.

The **VALUES** clause in the **INSERT** statements specifies the columns from the **SELECT** statement that have to be inserted into each tables.

Each row returned by the **SELECT** statement results in two insertions, one for the SAL_HISTORY table and one for the MGR_HISTORY table.

```
SQL> INSERT ALL
2      INTO sal_history          VALUES (empno,hiredate,sal)
3      INTO mgr_history          VALUES (empno,mgr,sal)
4      SELECT
5          empno, hiredate, sal, mgr
6      FROM      emp
7      WHERE     empno > 7400;

28 row(s) created.
```

6.1.3.Conditional INSERT ALL

This example is similar to the previous one.

The **INSERT** statement is referred to as a conditional **ALL INSERT**, as further restriction is applied to the rows that are retrieved by the **SELECT** statement. From the rows that are retrieved by the **SELECT** statement, only those rows in which the value of SAL column is more than 2000 are inserted into the SAL_HISTORY table, and similarly only those rows where the value of the MGR column is more than 7400 are inserted in the MGR_HISTORY table.

```
INSERT ALL
  WHEN sal > 2000 THEN
    INTO sal_history VALUES (empno,hiredate,sal)
  WHEN mgr > 7400 THEN
    INTO      mgr_history VALUES (empno,mgr,sal)
  SELECT     empno, mgr, hiredate, sal
  FROM      emp
  WHERE     empno > 7400

18 row(s) created.
```

6.1.4. Conditional FIRST INSERT

The following example inserts rows into more than one table, using one single **INSERT** statement.

The **SELECT** statement retrieves the details of department ID, total salary, and maximum hire date for every department in the EMP table.

The **INSERT** statement is referred to as a conditional **FIRST INSERT**, as an exception is made for the departments whose total salary is more than 1000 \$. The condition **WHEN SAL > 1000** is evaluated first. If the total salary for a department is more than 1000\$, then the record is inserted into the **SPECIAL_SAL** table irrespective of the hire date.

If this first **WHEN** clause is evaluated to true, the Oracle Server executes the corresponding **INTO** clause and skips subsequent **WHEN** clauses for this row.

For the rows that do not satisfy the first **WHEN** condition, the rest of the conditions are evaluated. If a row does not satisfy any **WHEN** condition the record is inserted into the **hiredate_history** table.

```
SQL> INSERT FIRST
2      WHEN SAL > 1000      THEN
3      INTO special_sal      VALUES (deptno,sal)
4      WHEN hiredate LIKE ('%00%') THEN
5      INTO hiredate_history_00      VALUES
(deptno,hiredate)
6      WHEN hiredate LIKE ('%99%') THEN
7      INTO hiredate_history_99      VALUES
(deptno,hiredate)
8      ELSE
9      INTO hiredate_history VALUES (deptno,hiredate)
10     SELECT      deptno,sum(sal)  sal,  MAX(hiredate)
hiredate
11     FROM      emp
12     GROUP BY      deptno;

4 row(s) created.
```

6.1.5. Pivoting INSERT

Pivoting is an operation in which you need to build a transformation so that each record from any input stream, such as, a nonrelational database table, must be converted into multiple records for more relational database table environment.

In the following statement, the sales data is received from the nonrelational database table **SALES_SOURCE_DATA**, which is the details of the sales performed by a sales representative on each day of a week, for a week with a particular week ID.

```
SQL> INSERT ALL
2      INTO sales_info VALUES (employee_id,week_id,sales_MON)
3      INTO sales_info VALUES (employee_id,week_id,sales_TUE)
4      INTO sales_info VALUES (employee_id,week_id,sales_WED)
5      INTO sales_info VALUES (employee_id,week_id,sales_THUR)
6      INTO sales_info VALUES (employee_id,week_id,sales_FRI)
7      SELECT employee_id, week_id, sales_MON, sales_TUE,
sales_WED
8      ,sales_THUR, sales_FRI
9      FROM sales_source_data;
```

```
5 row(s) created.
```

6.2.External tables

6.2.1.Creating an external table

An external table is a read-only table whose metadata is stored in the database but whose data is stored outside the database.

You can use external data as a virtual table. This data can be queried and joined directly and in parallel without requiring the external data to be first loaded in the database.

It's possible to use SQL, PL/SQL, and java to query the data in an external table.

The main difference between external and regular tables is that externally organized tables are read-only.

So, no DML operations are possible on it.

The metadata for an external table is created using the **CREATE TABLE** DDL statement.

The Oracle Server provides two major access drivers for external table.

One, the loader access driver, or **ORACLE_LOADER**, is used to read data from external files using the Oracle loader technology.

The other Oracle provided access driver, the import/export access driver, or **ORACLE_INTERNAL**, can be used for both the importing and exporting of data using a platform independent format.

You can create an external table using the **ORGANIZATION EXTERNAL** clause with the **CREATE TABLE** statement.

When you use it, you create a metadata in the data dictionary that you can use to access external data. If you specify **EXTERNAL** in the **ORGANIZATION** clause, you indicate that the table is read-only, and located outside the database.

TYPE access_driver_type indicates the access driver of the external table.

The **REJECT LIMIT** clause allows you to specify how many conversion errors can occur during a query of the external data before an Oracle error is returned and the query is aborted.

DEFAULT DIRECTORY lets you specify one or more default directory objects corresponding to directories on the file system where the external data sources may reside.

The **ACCESS PARAMETERS** clause allows you to assign values to parameters of the specific access drivers for this external table.

LOCATION is a clause that lets you specify one external locator for each external data source. Usually, the **location_specifier** is a file, but it need not be.

6.2.2.Example of creating an external table

You have to use the **CREATE DIRECTORY** statement to create a directory object, which is an alias for a directory on the server's file system where an external data source resides.

Note that you must have the **CREATE ANY DIRECTORY** system privileges to create directories. Once you created a directory, you automatically grant the **READ** object privilege and can grant **READ** privileges to other users and roles.

Syntax:

CREATE [OR REPLACE] DIRECTORY *directory* **AS** '*path_name*';

OR REPLACE: Re-create the directory database object if it already exists.

Directory: Specify the name of the directory to create.

Path_name: Specify the full pathname of the operating system.

Assume that there is a flat file that has records in the following format:

```
10, Jones, 11-DEC-1934
20, Smith, 12-JUN-1972
```

Records are delimited by new lines, and the fields are terminated by a comma. The name of the file is: /flat_files/empl.txt

To convert this file as the data source for an external table, whose metadata will reside in the database, you must:

1. Create a directory object emp_dir as follows:

```
CREATE DIRECTORY emp_dir AS '/flat_files';
```

2. Run the following **CREATE TABLE** statement.

```
SQL> CREATE TABLE      odlemp (
 2      empno NUMBER, ename CHAR (20), hiredate DATE)
 3      ORGANIZATION EXTERNAL
 4      (TYPE ORACLE_LOADER
 5      DEFAULT DIRECTORY emp_dir
 6      ACCESS PARAMETERS
 7      (RECORDS DELIMITED BY NEWLINE
 8      BADFILE          'bad_emp'
 9      LOGFILE           'log_emp'
10      FIELDS TERMINATED BY ','
11      (empno CHAR,
12      ename  CHAR,
13      hiredate CHAR date_format date mask "dd-mon-yyyy"))
14      LOCATION          ('empl.txt'))
15      PARALLEL          5
16      REJECT LIMIT      200;

Table created.
```


-> The **TYPE** specification is given to illustrate its use.

ORACLE_LOADER is the default access driver if not specified.

ACCESS PARAMETERS provides values to parameters of the specific access driver.

The **PARALLEL** clause enables five parallel execution servers to simultaneously scan the external table when executing **INSERT INTO TABLE** statement.

6.2.3. Querying external tables



An external table does not describe any data that is stored in the database.

It describes how the external table layer needs to present the data to the server.

When the database needs to access data in an external source, it calls the appropriate access driver to get the data from an external source in a form that the database server expects.

Note that the description of the data in the data source is separate from the definition of the external table.

The data types for the fields in the data source can be different from the columns in the table.

The access drivers ensure the data from the data source is processed so that it matches the definition of the external table.

6.3. CREATE INDEX with CREATE TABLE statement

It's possible for a user who has sufficient privileges to create an index while creating a table.

In the following example, the **CREATE INDEX** clause is used with the **CREATE TABLE** statement to create a primary key index explicitly.

You can name the index at the time of **PRIMARY KEY** creation.

```
SQL> CREATE TABLE new_emp
2      (empno NUMBER(6)
3          PRIMARY KEY USING INDEX
4          (CREATE INDEX emp_id_idx ON
5            new_emp(empno)),
6      ename VARCHAR2(50));

Table created.
```

You can ensure that the index has been created by querying the USER_INDEX data dictionary table.

```
SQL> SELECT index_name, table_name  
2   FROM user_indexes  
3   WHERE table_name='NEW_EMP';
```

INDEX_NAME	TABLE_NAME
EMP_ID_IDX	NEW_EMP